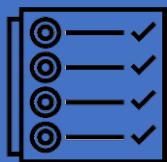




Implement Relational Data Stores

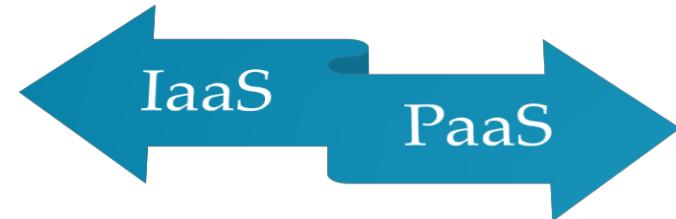
Implement relational data stores



Learning Objectives

- Azure SQL Database offerings
 - PaaS vs IaaS
 - PaaS deployment options
 - Purchasing models and Service Tier
 - Elastic Database
 - Managed Database
 - Security
- Azure SQL Datawarehouse
 - Traditional vs Modern Architecture
 - Azure Synapse Analytics
 - MPP Architecture
 - Storage and Sharding Pattern
 - Data Distribution and Table types
 - Partitioning
- Loading methods
 - PolyBase vs SSIS
 - PolyBase steps and demos
- Azure SQL Server vs Datawarehouse

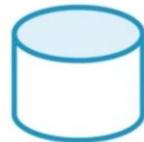
Introduction



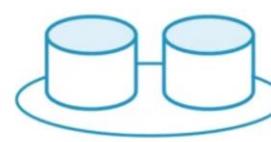
Azure SQL Database Deployment Types

PaaS

RELATIONAL - DATABASE - AS - A - SERVICE



Single Database
Own set of resources



Elastic Pool
Collection of databases sharing resources



Managed Instance
Dedicated engine instance running collection of databases

SQL SERVER
INSTANCE

Why SQL Server in Azure?



Fully Managed



Predictable
performance
and pricing



DISASTER RECOVERY

Geo-replication
and restore
services



DB COLLECTION

Elastic pool for
unpredictable
workloads



DTU COMPUTE
SHARED RESOURCES

Supports existing SQL
Server tools, libraries,
and APIs



Scalability with no
downtime



99.99%
availability
built-in



Secure and compliant
for your sensitive data

\$2M DOWN
PER YEAR

Azure IaaS vs PaaS Database offerings?



SQL Server on Azure VMs
SQL Server inside a
fully-managed VM in Azure

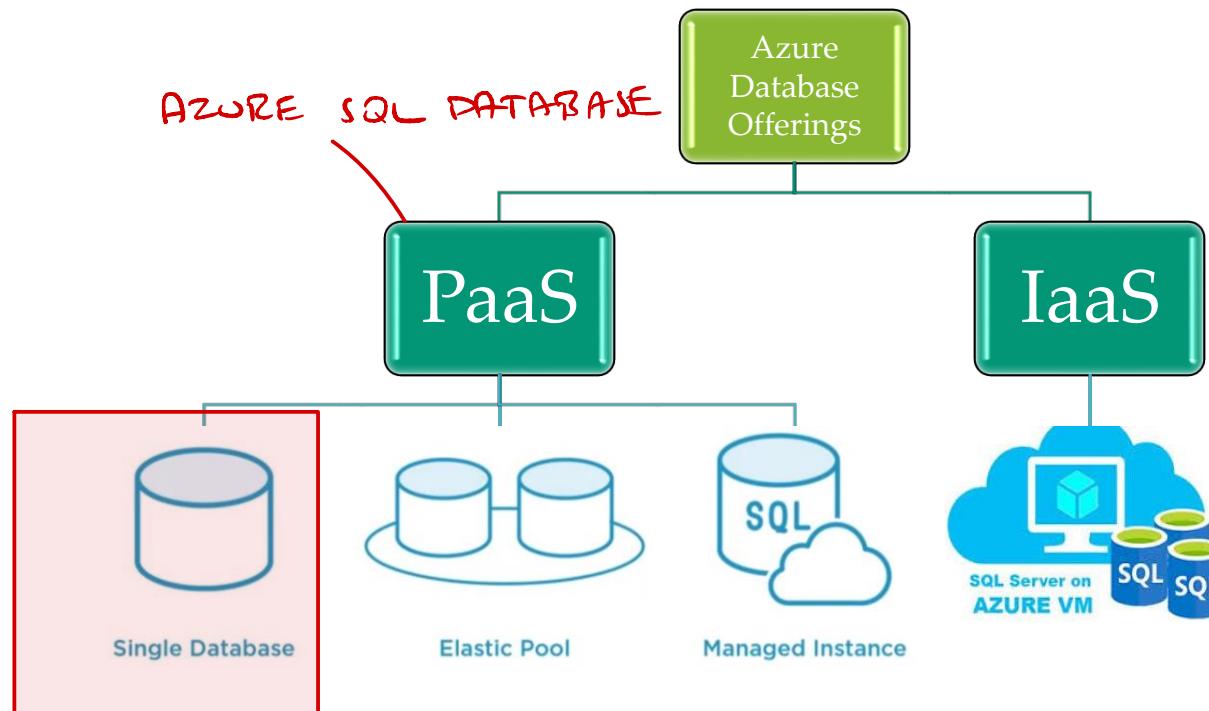


Azure SQL Database
Database-as-a-service (DBaaS)
hosted in Azure

SQL Server and Azure SQL Database Comparison

SQL Server	Azure SQL Database
Responsibility Comparison	
Maintain operating system	Choosing the service tier
Maintain SQL Server	Testing HA & DR
Maintain HA & DR	Maintain performance tuning
Maintain security	Maintain change control
Maintain performance tuning	Maintain security
Maintain change control	
Benefit Comparison	
99.95 % availability	99.99% availability
Full control over SQL Server engine	Change resources with no downtime
Full parity with matching version of on-premises SQL server	Most commonly used SQL Server features available
Easy migration from SQL Server on-premises	Built-in backups, patching, recovery
Private IP address within Azure VNet	Ability to assign necessary resources to individual databases (CPU, Storage)
	Built-in advanced intelligence and security

Azure Database Deployment options



SQL SERVER IS A LOGICAL SERVER (MANAGED)

SQL Server(PaaS) Deployment Options



Single database

Each DB with its own
guaranteed compute,
memory, and storage

ISOLATED DB
DTU OR VCORES



Elastic pool

Fixed resources will be
shared by all databases
in the pool

FIXED RESOURCES
ARE DYNAMICALLY
SHARED



Managed instance

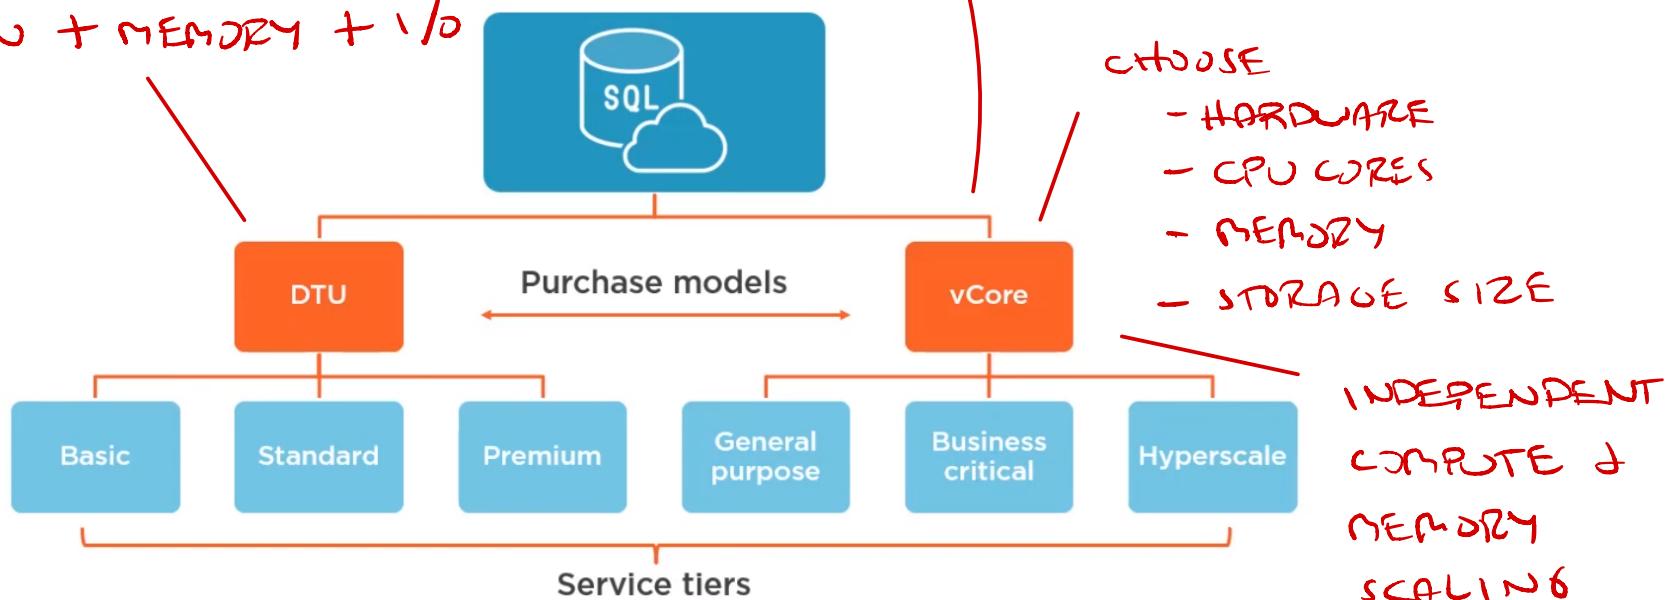
Each managed
instance has its
guaranteed resources

- SET OF DB'S THAT CAN BE USED TOGETHER
- . EASY MIGRATION

DEDICATED
SQL SERVER
INSTANCE
MANAGED ON
VM

Azure SQL Database Service Tiers

CPU + MEMORY + I/O



vCore-based vs DTU-based Model

vCore-based

ALL DEPLOYMENT OPTIONS

- For Single database, elastic pool and managed instance
- Best for customers who need flexibility, control, and transparency
- Straightforward way to translate on-premises workload to the cloud
- Microsoft recommends vCore-based model

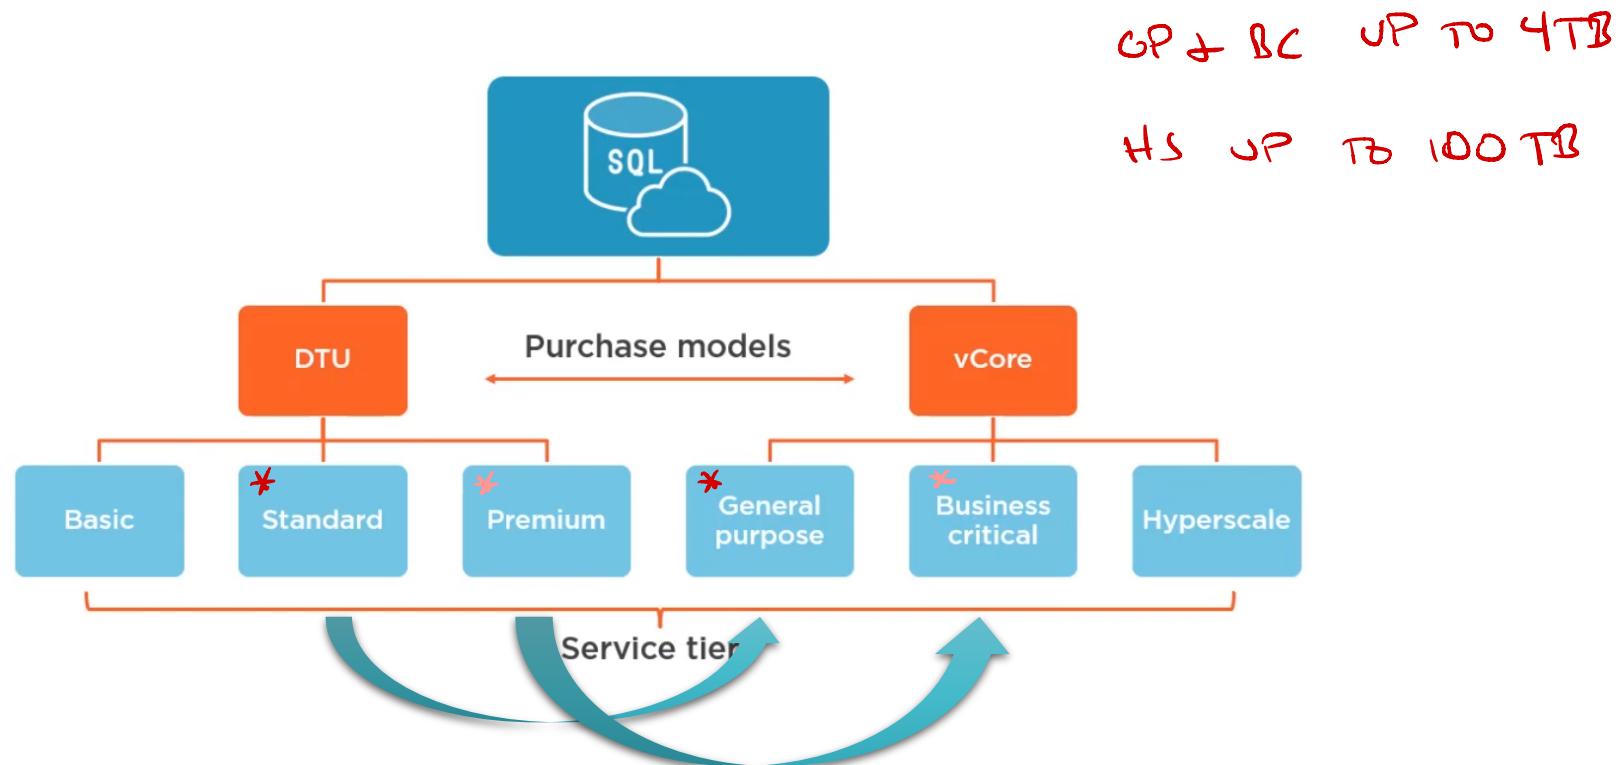
DTU-based

- Only for single database and elastic pool
- Best for customers who want single, preconfigured resource options
- Might need to calculate the needed DTUs before migration
- If the DTU-based purchasing model needs your performance and business requirements, you should continue using it.

Converting DTU-based Model to vCore-based

- If your single database or elastic pool consumes more than 300 DTUs, converting to the vCore-based model might reduce your costs.
- You can use API of your choice or Azure Portal to convert to vCore based model with 'no downtime'.
- Azure SQL Database managed instance only supports vCore-based purchasing model.

Azure SQL Database Service Tiers



Azure SQL Database Options

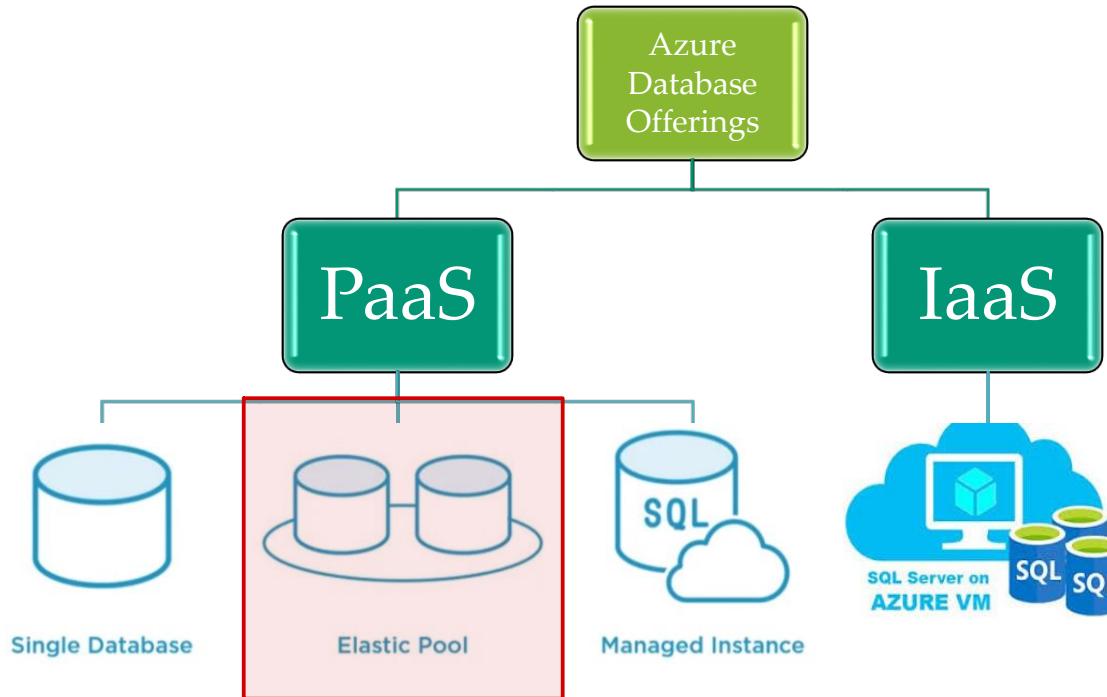
SINGLE DB
ELASTIC POOL
MANAGED INST

DTU
VCORE

DTU B / S / P
VCORE GP / BC / HS



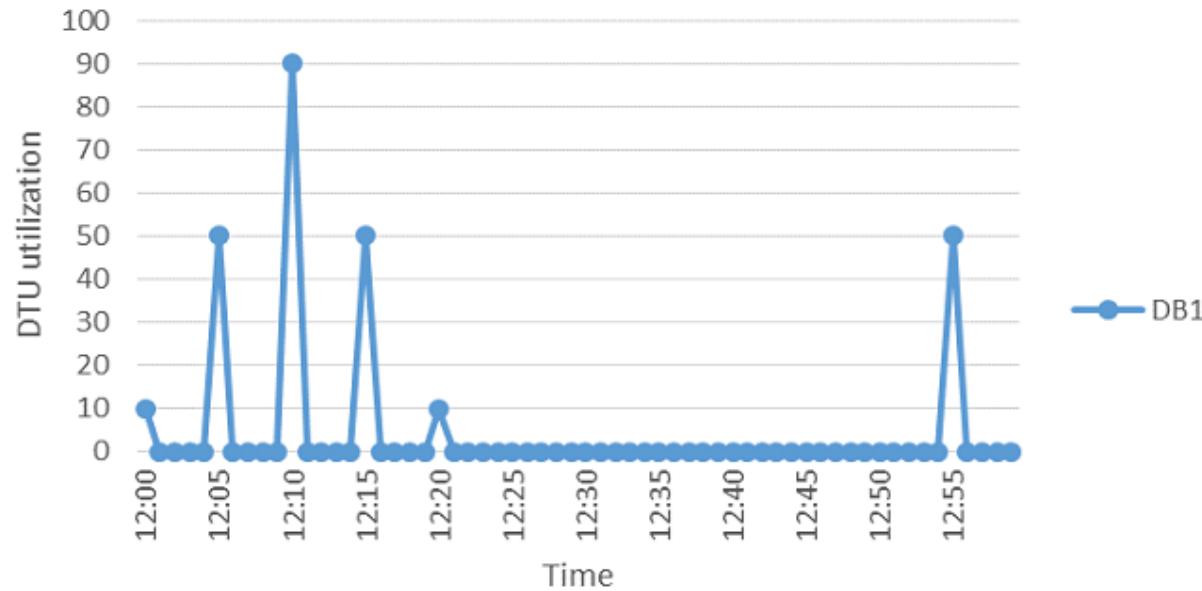
Azure Database Elastic Pool



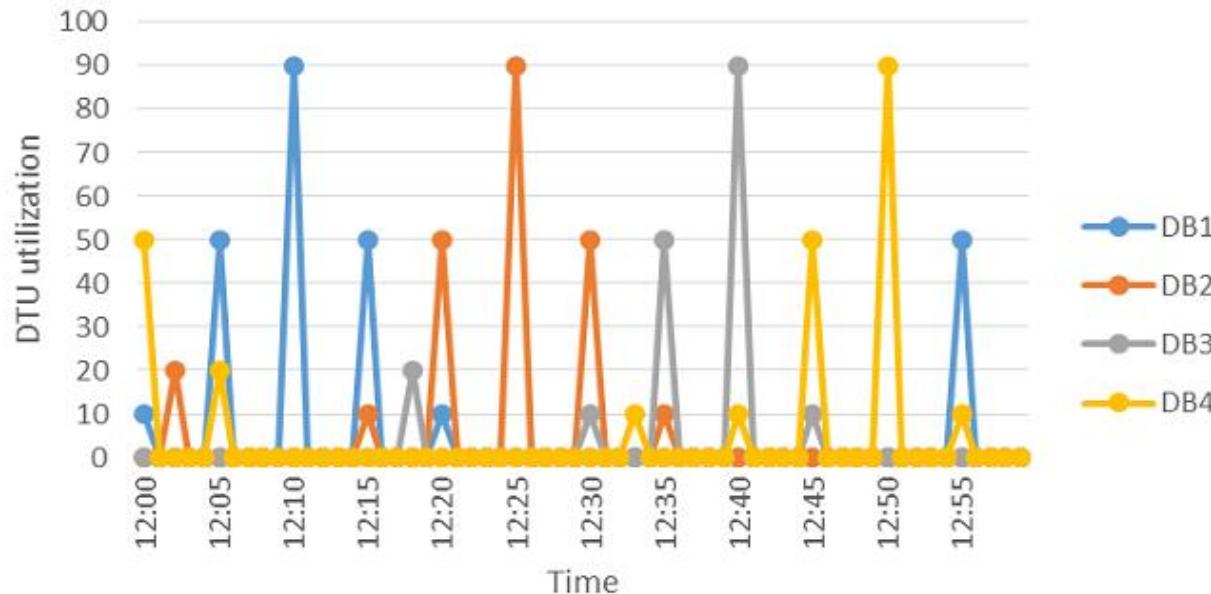
SQL Database Elastic Pools

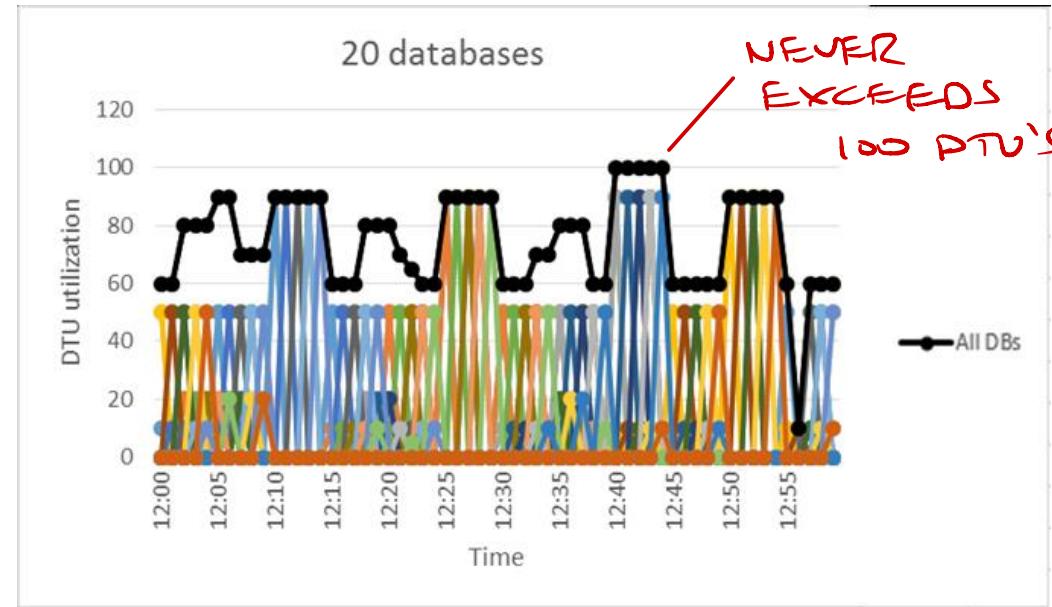
An Azure SQL Elastic Pool allows you to allocate a shared set of compute resources to a collection of Azure SQL databases,

1 database



4 databases

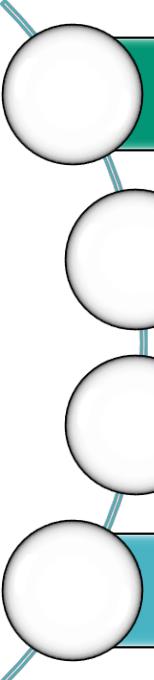




ALL 20 DB'S
CAN SHARE
THE 100 DTU'S
INSTEAD OF
CONSUMING
OWN RESOURCE

SHARED DTU'S FOR ALL DATABASES

Azure SQL Database Elastic Pool



So Azure SQL elastic pool is a cost- effective solution for managing and scaling multiple databases that have varying and unpredictable usage demands.

The databases in an elastic pool are on a single Azure SQL Database server and share a set number of resources at a set price.

Elastic pool enables developers to optimize the price performance for a group of databases within a prescribed budget. *

Elastic pools prevent over-provisioning or under-provisioning of resources.

Security

4 LAYERS OF DEFENSE



Network security

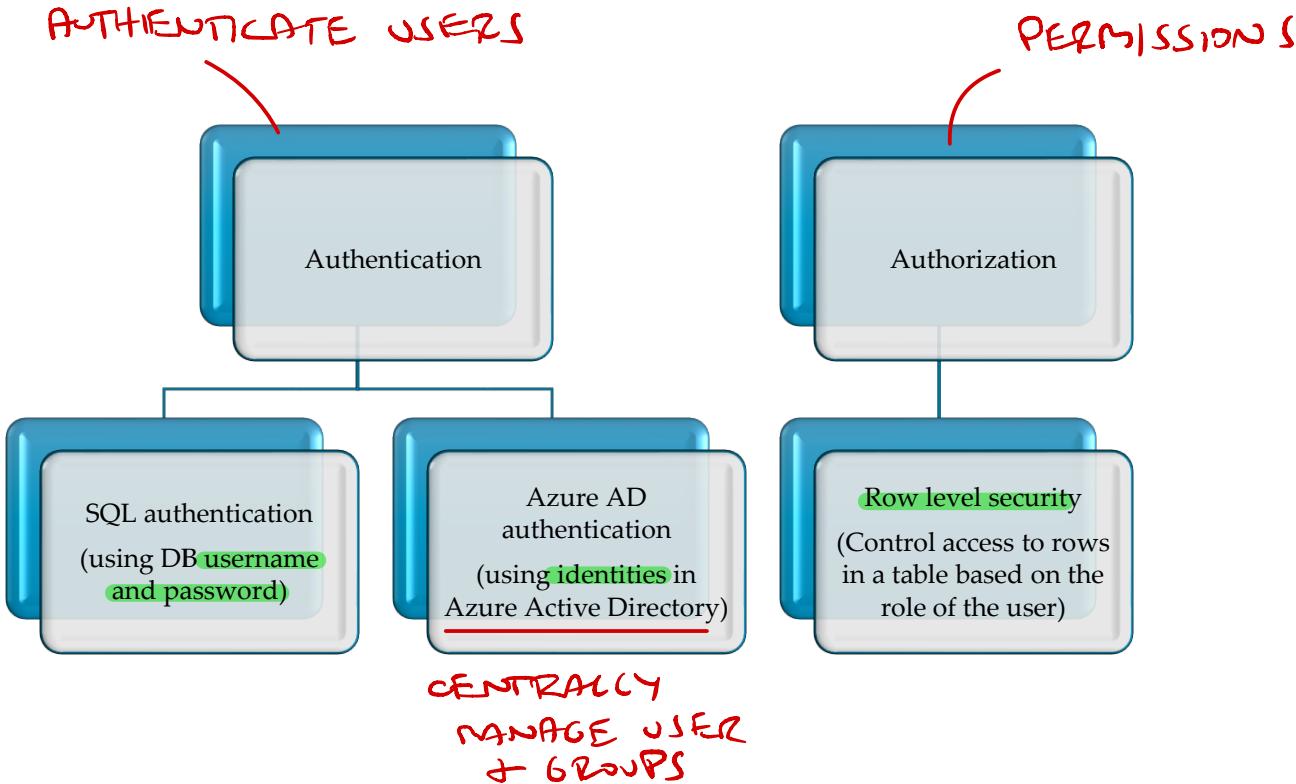
IP firewall rules

- Grant access to databases based on the originating IP address of each request.

Virtual network firewall rules

- Enable Azure SQL Database to only accept requests originating from subnets inside a virtual network.

Access Management



Threat Protection

SQL AUDITING

in Azure
Monitor logs
and Event
Hubs

Advanced
Threat
Protection

- Tracks database activities and helps maintaining compliance with security standards

- Analyzes your SQL Server logs to detect unusual behavior and potentially harmful attempts

SQL INJ
ROUTE
FORGE
ETC.

Information Protection

ALL DATA IN TRANSIT IS ENCRYPTED (TLS)

Transport Layer Security TLS

- Always enforces encryption for all connections

Transparent Data Encryption TDE

- (Protects data at rest from offline access to raw files or backups)

DEFAULT AES ENCR
KEYS IN AZURE KEY VAULT

Dynamic Data masking DDM

- (Protects sensitive data by masking it for non-privileged users)

Security Management

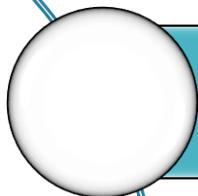
Vulnerability assessment

- Discover track and remediate potential **database vulnerabilities**.

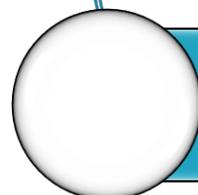
Data discovery & Classification

- Identify and **label sensitive data** for monitoring and alerting

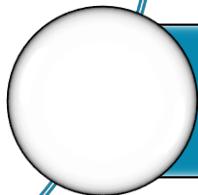
Managed Instance Advance Security



Native virtual network implementation and connectivity to your on-premises environment using Azure Express Route or VPN Gateway.



In a default deployment, SQL endpoint is exposed only through a private IP address, allowing safe connectivity from private Azure or hybrid networks.



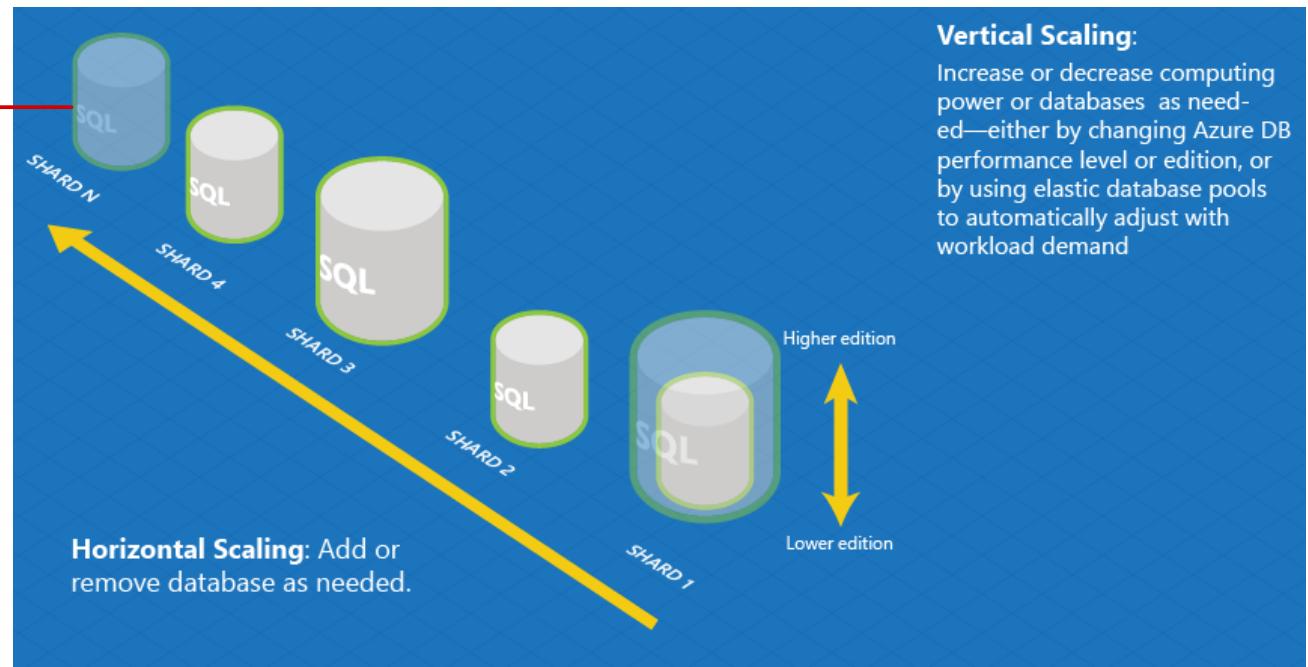
Single-tenant with dedicated underlying infrastructure (compute, storage).

Azure SQL Database Scaling

Azure SQL Database supports two types of scaling:

- **Vertical scaling:** Scale up or down the database by adding more compute power.
- **Horizontal scaling:** Scale out or Add more databases and to shard your data into multiple database nodes.

EACH SHARD DB CAN
BE SCALING
INDEPENDENTLY



Azure SQL Database Scaling

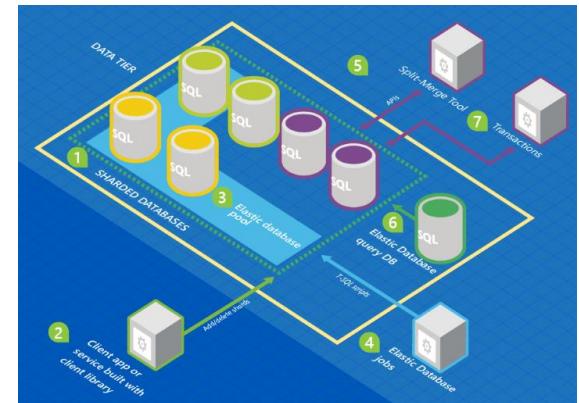
HORIZONTAL SCALING OPTIONS

Read Scale Out

- Allows you to use the capacity of the read-only replicas for read-only queries.
- Feature is intended for the applications that include logically separated read-only workloads, such as analytics
- Benefit: When Secondary nodes handles heavy reports and analytical queries, primary writable node saved resources that might be used to improve the performance
- Secondary nodes is asynchronous POSSIBLE SMALL DELAY
- Premium (DTU-based model) or in the Business Critical (vCore-based model)
- Also available in Hyperscale with secondary replica creation

Global Scale-out/Sharding

- Split your data into multiple database nodes. / NOT COPIES
- Every database shard is an independent database where you can add or remove resources as needed.
- Application may access only the shard associated to that region without affecting other shards.
- Why?
 - Data or transaction throughput exceed the capabilities of individual database
 - Tenants may require physical isolation
 - Different sections of a database may need to reside in different geographies for compliance, performance, or geopolitical reasons.



EXAMPLE: US SHARD , EU SHARD (ORG DOMAIN SHARDS)

Vertical vs Horizontal Scaling

Azure SQL Database supports two types of scaling:

- Vertical scaling:
 - Scale up or down the database by adding more compute power.
OR REDUCING
 - CPU Power, Memory, IO throughput, and storage
 - DTU and vCore models to scale
 - Dynamic Scalability (Note: this is not auto-scale)
*MANUALLY ↑ ↓
RESOURCES*
 - Any change that you made will be almost instant.
- Horizontal scaling: Scale out or Add more databases and to shard your data
into multiple database nodes.
*EACH SHARD CAN BE
VERTICALLY SCALLED*



Azure SQL Database Scaling

VERTICAL SCALING OPTION

Change Service Tier

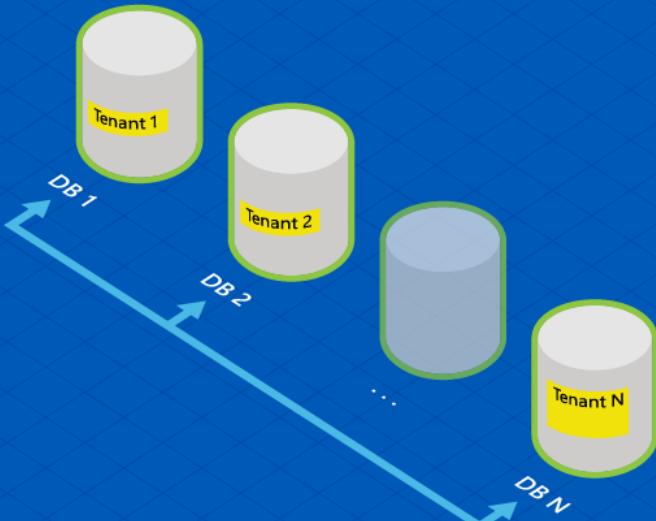
- From Standard/General Purpose to Premium/Business Critical.
- In Standard/General Purpose - Data stored on Azure premium disks
- In Premium/Business Critical - Data stored on local SSD



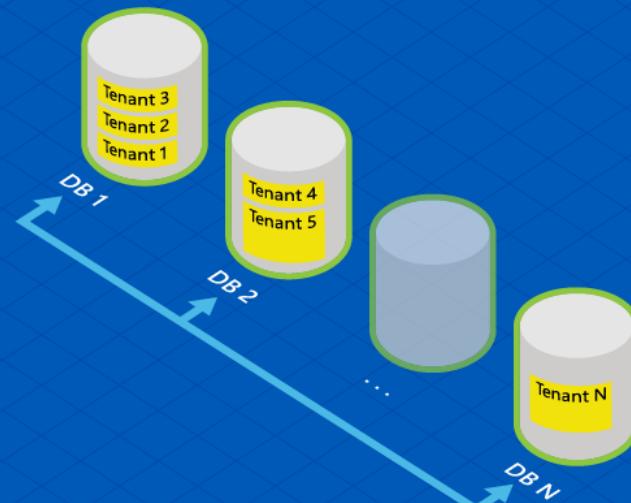
Single Tenancy vs Multi Tenancy

TENANT = CUSTOMER

Single Tenancy

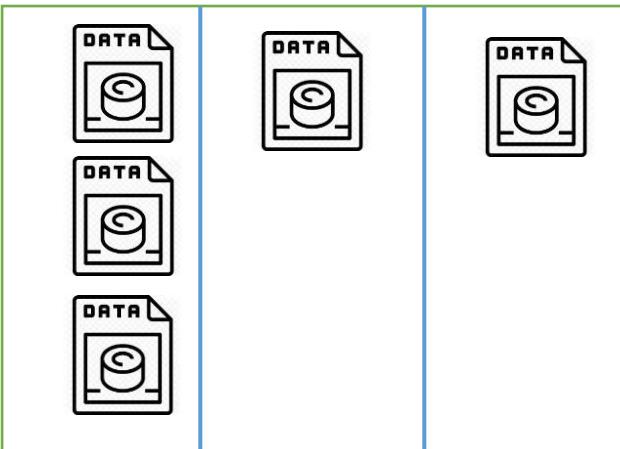


Multi Tenancy

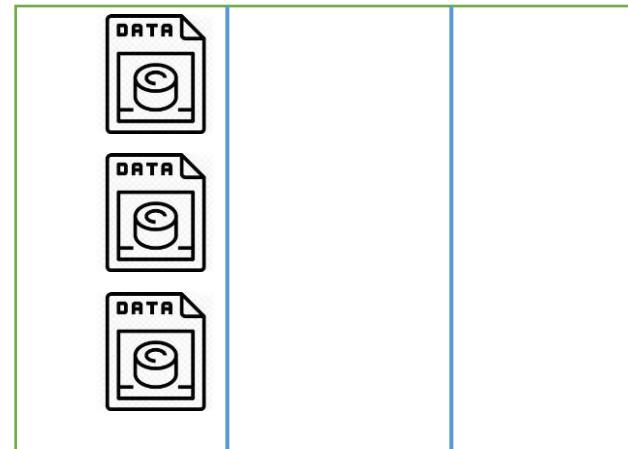


Azure SQL DB – HA and DR Options

Region A



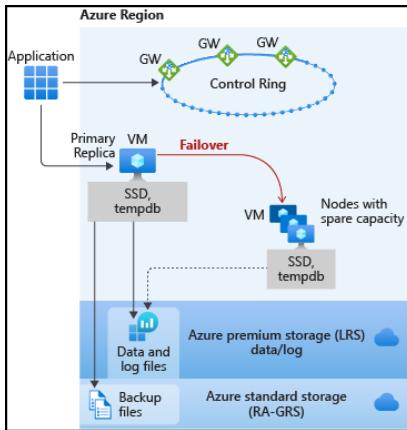
Region B (Read)



Each cluster is physically separate in what's called an availability zone, with its own separate utilities and networking.

Hundreds of miles away from the primary region to prevent data loss in the event of a natural disaster.

Azure SQL DB – High Availability Architecture

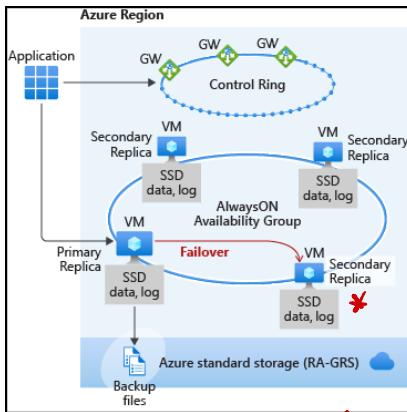


Standard availability model

- Separation of compute and storage
- Performance degradation during maintenance activities
- Budget-oriented business applications
- Basic, Standard, and General Purpose service tier availability

FAILOVERS CAN BE PLANNED OR UNPLANNED

SPARE TAKE OVER FROM PRIMARY NODE DURING FAILOVER.



Premium availability model

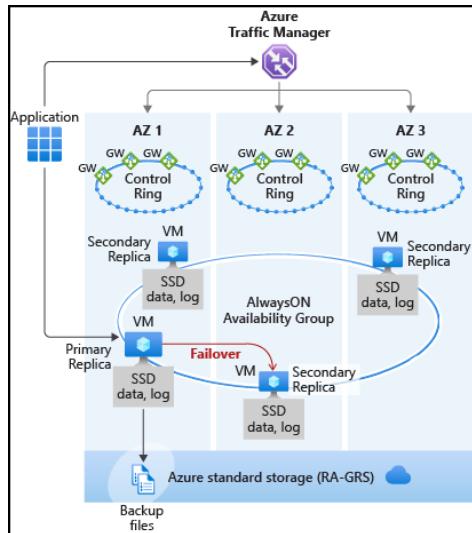
- Integrates compute resources and storage (locally attached SSD) on a single node
- Replicating both compute and storage to additional nodes creating a three to four-node cluster.
- Data is synchronized to at least one secondary replica before committing each transaction.
- Targets mission critical applications with high IO performance and high transaction rate
- Guarantees minimal performance impact during maintenance activities.
- Premium and Business Critical service tier availability
- Read Scale-Out feature in configure screen

* CAN BE CONFIGURED AS A READ-ONLY NODE FOR FAILURES *

Azure SQL DB –Zone Redundancy

AZ = AVAILABILITY ZONES

SETTING IN CONFIGURE SCREEN



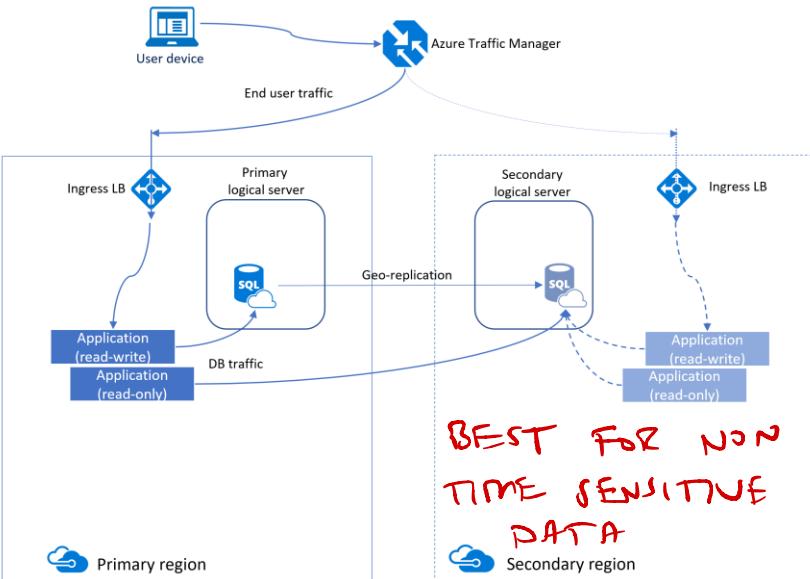
Zone redundant configuration

(SIMILAR TO PREMIUM AUR MODEL)

- Place different replicas of the Business Critical database to different availability zones in the same region
- Does not create additional database redundancy
- Enable it at no extra cost
- Supported in the Premium and Business Critical service tiers
- Not available in SQL Managed Instance.
- Increased network latency may increase the commit time and thus impact the performance of some OLTP workloads.

B/C DATA IS SYNCHRONIZED TO
SECONDARY REPLICA BEFORE
TRANSACTION IS COMMITTED

Azure SQL DB – Geo Replication



BEST FOR NON
TIME SENSITIVE
DATA

THERE WILL BE
A LAG IN REPLICATION

- Issues:
 - Supports only manual failover
 - End-point connection must be changed in the application after the failover
 - Must have the same firewall rules and the logins to run applications successfully without any discrepancies
- Create a readable secondary database in the same region or cross-region
- Use cases:
 - Can failover to the secondary database in case of an outage
 - Migrate a database from one server to another server in the same or cross region with minimal downtime.
- We can create up to four secondaries for each primary database.
- Data Loss:
 - Uses the Always-on feature to replicate committed transactions to the secondary database asynchronously.
 - May lag the primary database at any point in time.
- Manual - Forced Failover
 - This will make your secondary database immediately online and start accepting connections. **Forced failover may result in data loss.**

CAN BE MITIGATED
w/ FAILOVER GROUPS

GROUPS OF DATABASES CAN BE ADDED TO FAILOVER GRPS.

Compare geo-replication with failover groups *

Auto-failover groups simplify the deployment and usage of geo-replication and add the additional capabilities as described in the following table:

	Geo-replication	Failover groups
Automatic failover	No	Yes
Fail over multiple databases simultaneously	No	Yes
User must update connection string after failover	Yes	No
SQL Managed Instance support	No	Yes
Can be in same region as primary	Yes	No
Multiple replicas	Yes	No
Supports read-scale	Yes	Yes

LISTENER ENDPOINT IN FAILOVER GROUPS DO NOT CHANGE

AZURE SQL DB

BACKUP & RESTORE

TYPES OF BACKUPS

1

Full Backup

- Backs up the whole database
- Taken every week

2

Differential backup

- Captures only the data that has changed since the last full backup
- Taken every 12 hours

Backup
Frequency

3

Transaction log backup

- Records of all the committed and uncommitted transactions
- Taken every 5-10 mins

CONFIGURE BACKUPS IN
SERVER SETTINGS IN AZURE

Storage cost and Security of Backup files



Storage Cost

Backup files are copied to RA-GRS standard blob storage by default to paired region



Security

Backup are automatically encrypted at rest using TDE, blob storage is also protected

BACKED UP
TWO REGIONS. SECONDARY IS READ-ONLY

Backup Retention Period



Backup storage redundancy

- Point in time restore (PITR) - 7-35 days
- Long-term retention - Up to 10 years
→ LTR

Long term retention (LTR)

- One or more long term retention periods to your database to meet regulatory, compliance or other business purposes
- Full backups can be taken up to 10 years
- Stored in RA-GRS blob storage
- Any change of the LTR policy applies to the future backups

LTR and Managed Instance

- LTR is not yet available for databases in Managed Instances
- You can use SQL Agent jobs to schedule copy only database backups as an alternative to LTR beyond 35 days
- These backups can be kept in the Azure blob storage

Backup Restore



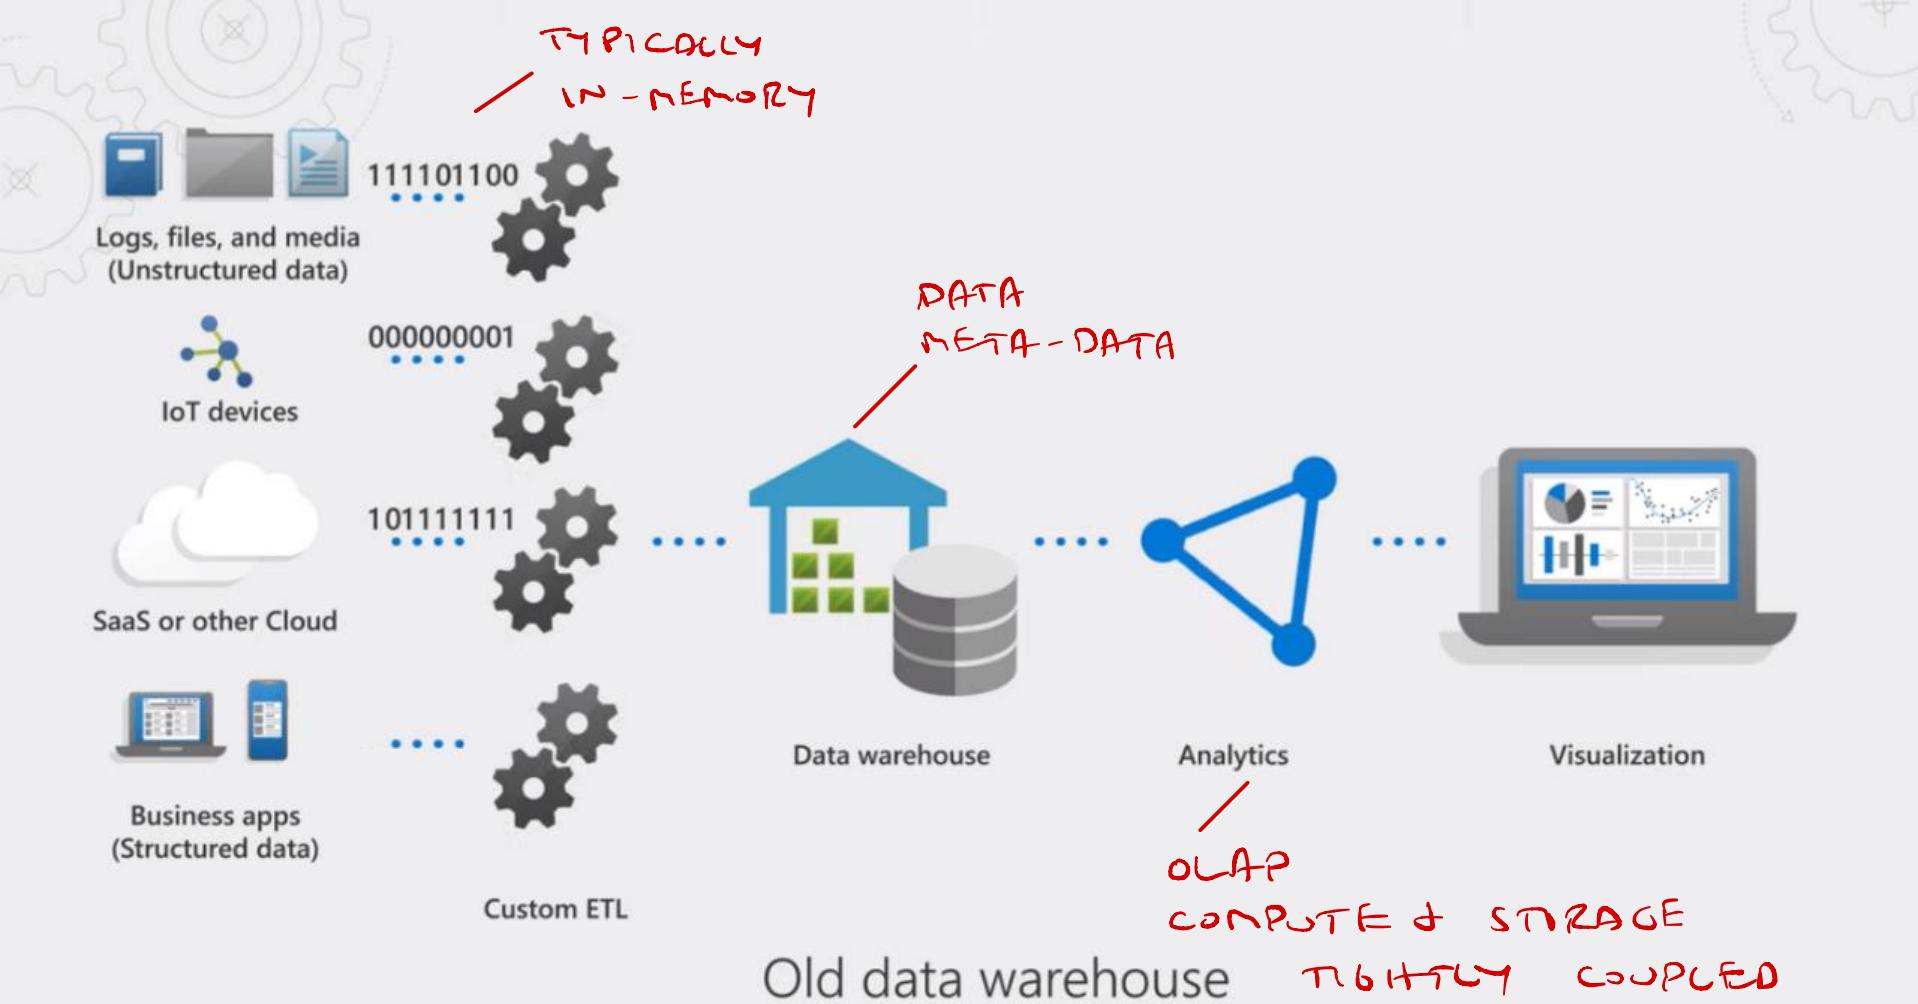
Backup usage

- Point-in-time restore of existing database
- Point-in-time restore of deleted database
- Geo-restore
- Restore from long-term backup

• *If you delete an Azure Logical SQL Server, all elastic pools and databases that belong to that logical server are also deleted and cannot be restored*

Restore Time is impacted By

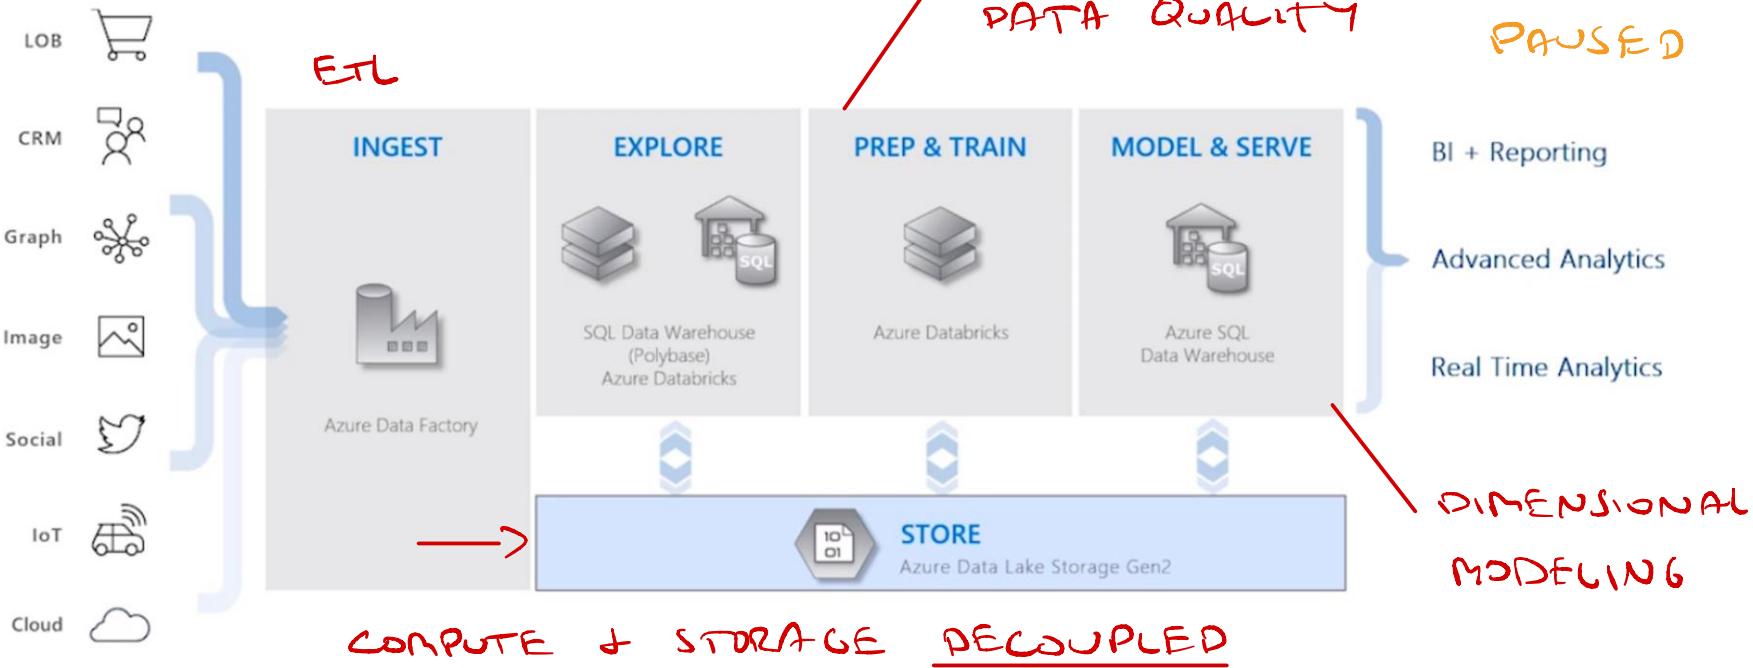
- Size of database & Compute size of the database
- Number of transaction logs and Amount of activity
- Network bandwidth if the restore is to a different region
- Concurrent restore requests being processed region



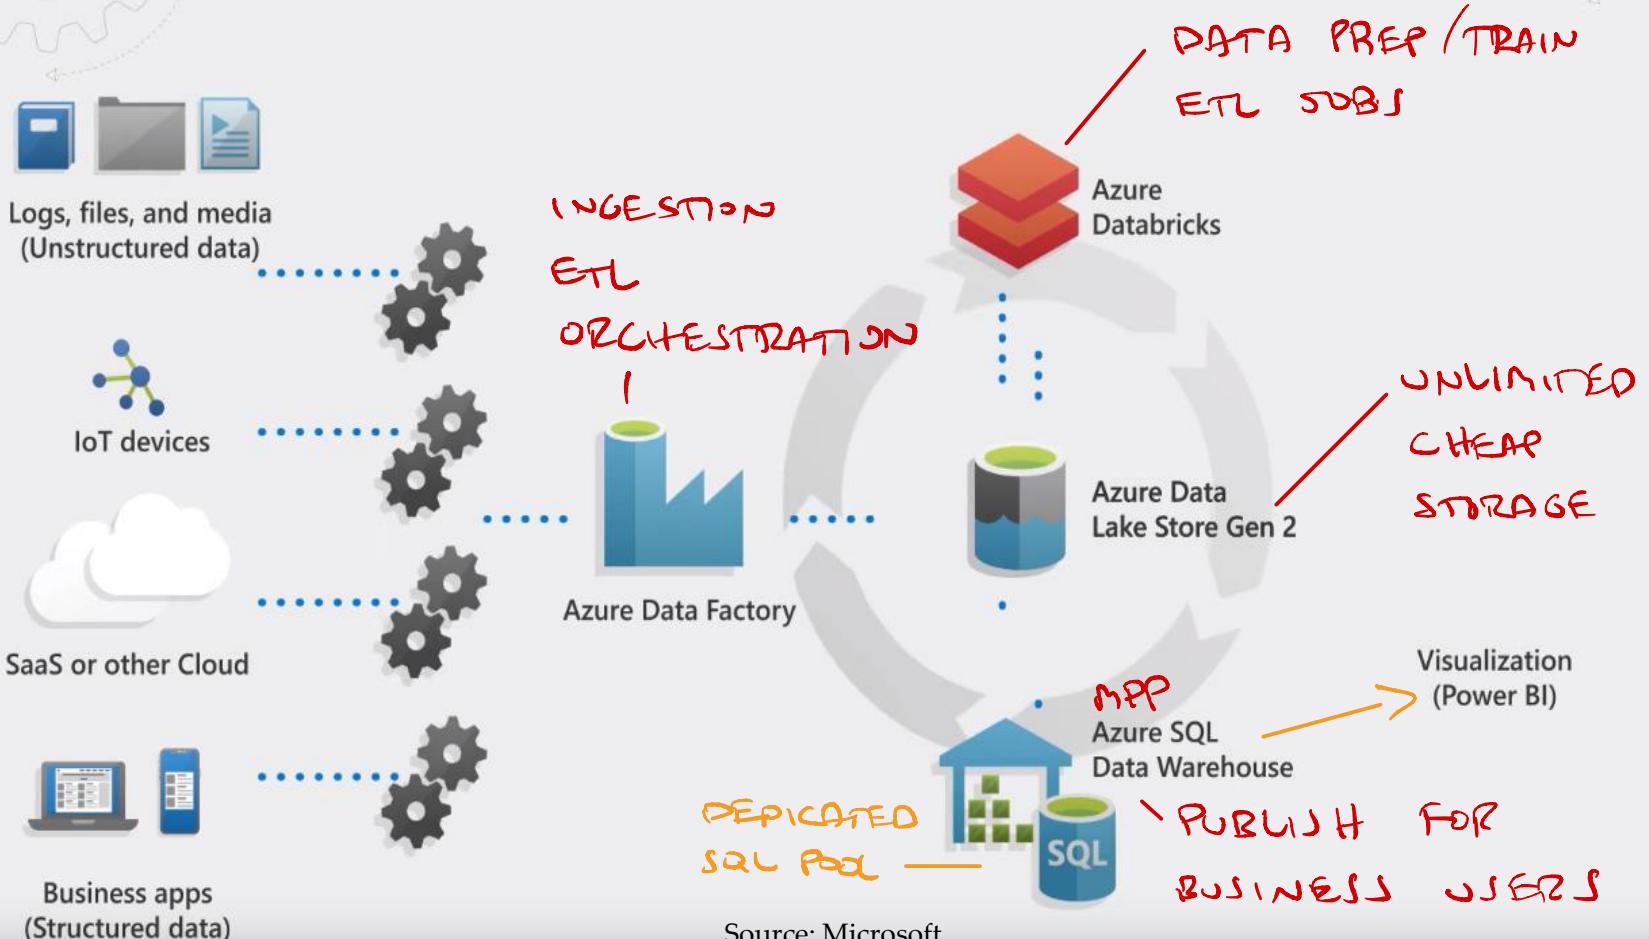
COMPUTE DWU = CPU + MEMORY + IO

• DATA WAREHOUSING UNIT

Modern Data Warehousing



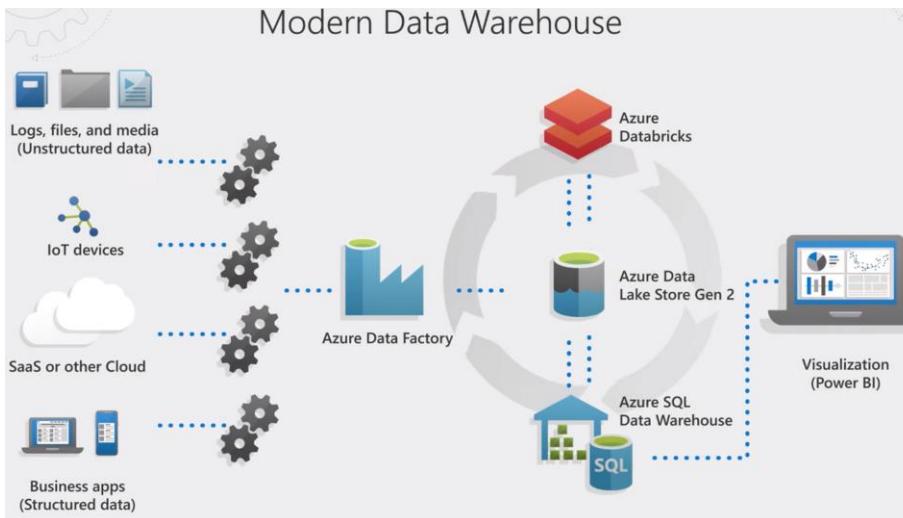
Modern Data Warehouse IN TERMS OF SERVICES



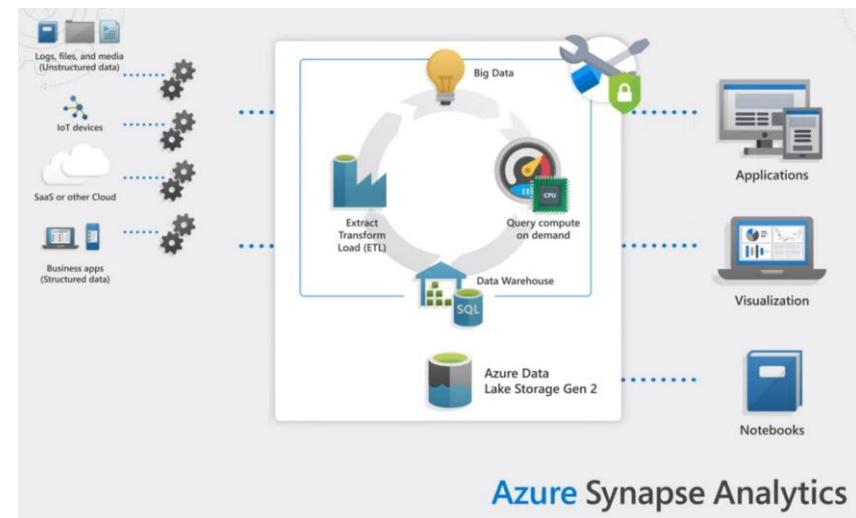
Azure Synapse Analytics

"Synapse is the next generation of Azure SQL Data Warehouse, blending big data analytics, data warehousing, and data integration into a single unified service that provides end-to-end analytics with limitless scale."

Modern vs Synapse Architecture

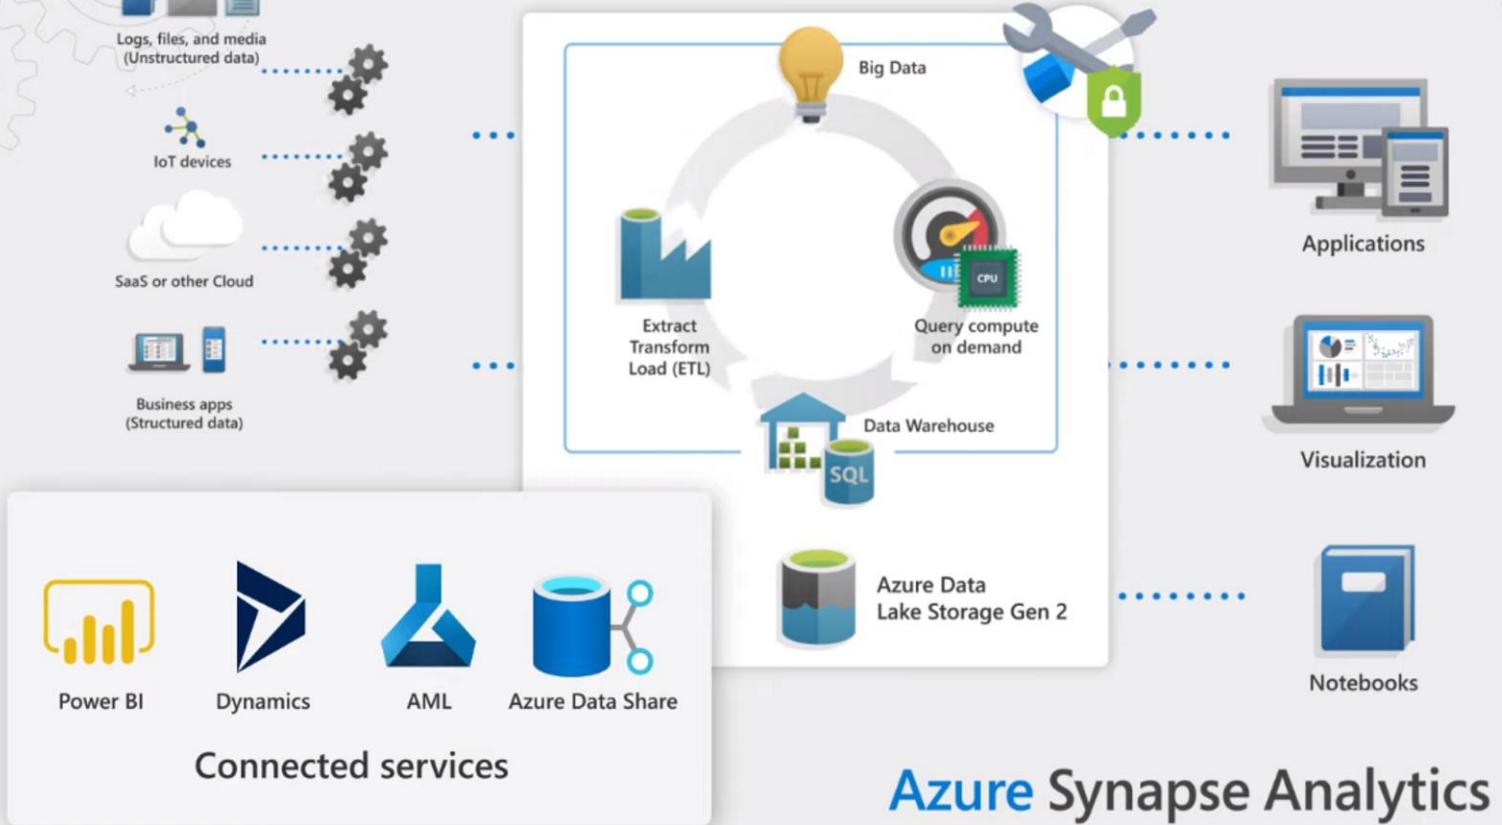
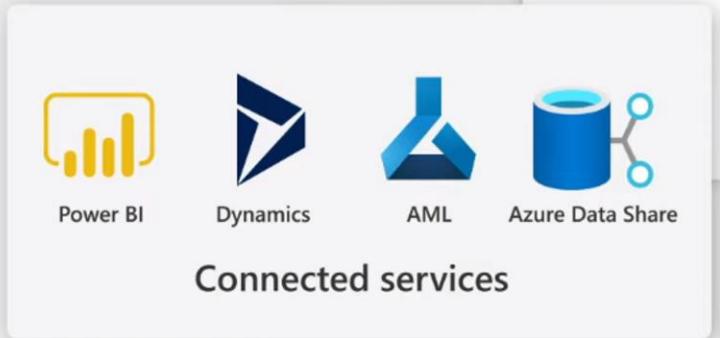


Modern Data Warehouse Architecture

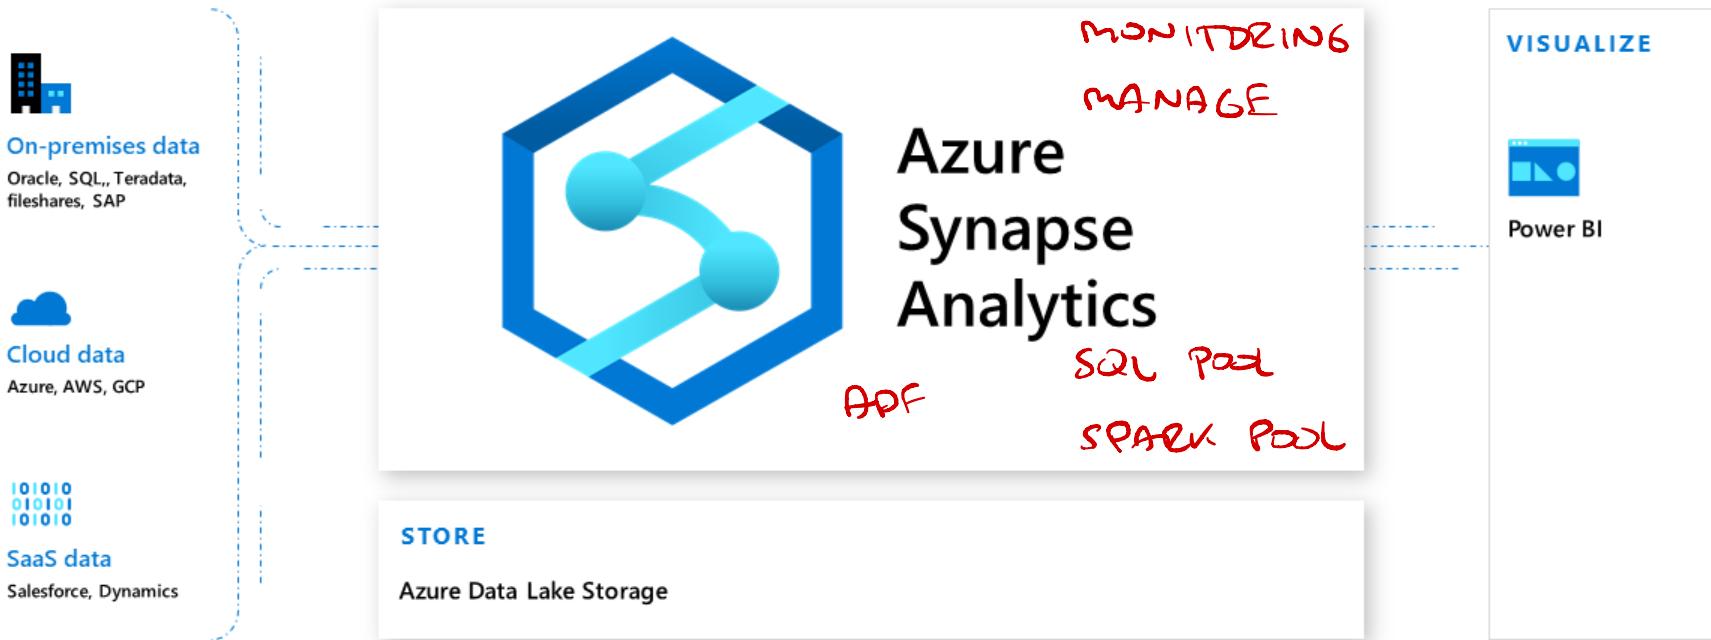


Azure Synapse Analytics

Synapse Analytics Architecture

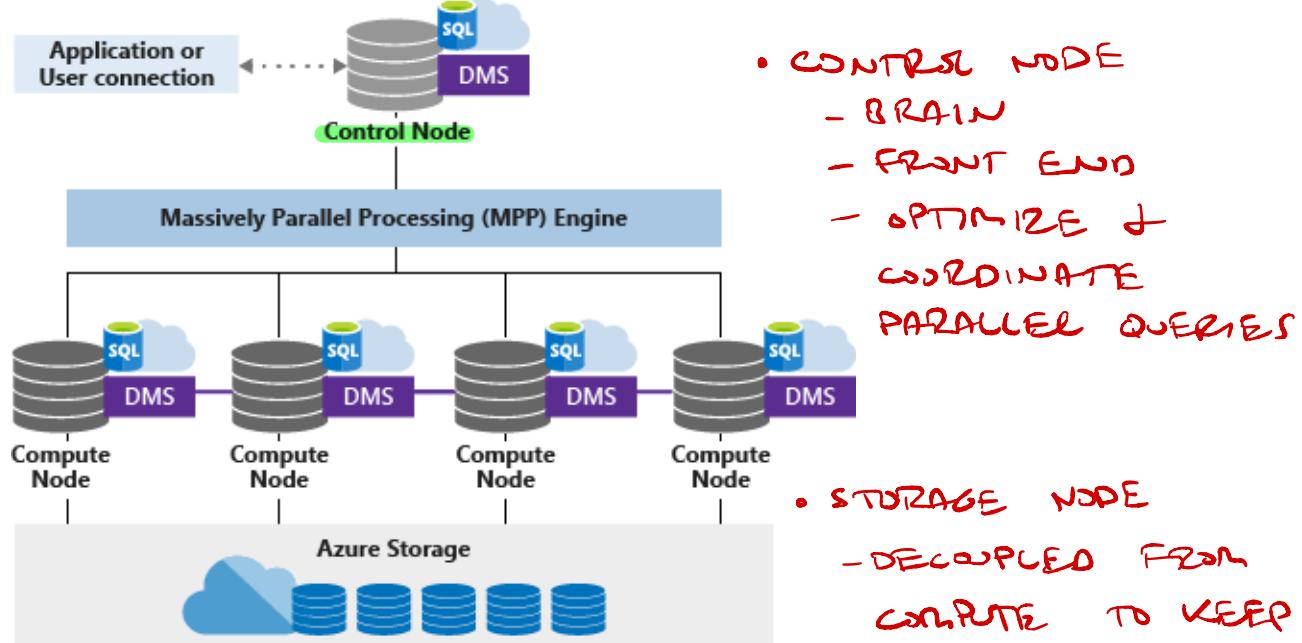


Azure Synapse Analytics



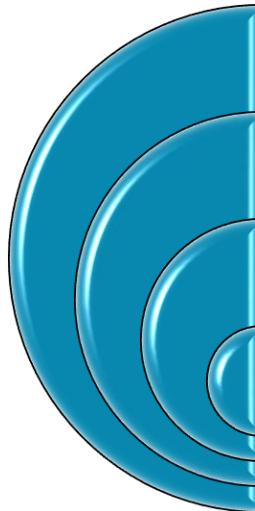
Azure Synapse MPP Architecture

NODE BASED ARCHITECTURE



DWU	Loading 3 Tables	Ran Report
100	15	20
500	3	4

Azure Storage and Distribution



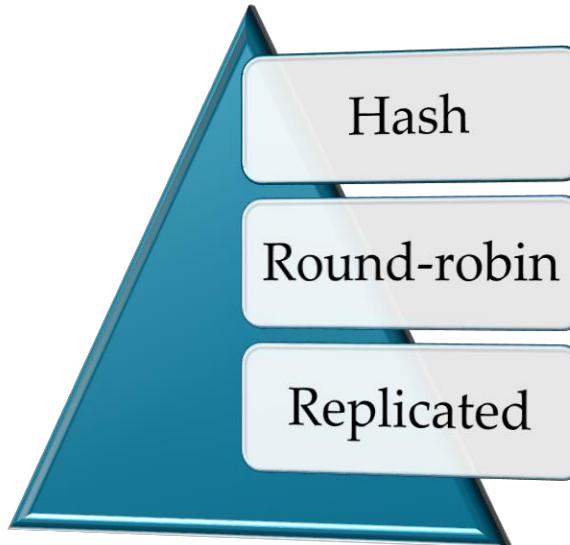
SQL DW charges separately for storage consumption

A distribution is the basic unit of storage and processing for parallel queries

Rows are stored across 60 distributions which are run in parallel

Each compute node manages one or more of the 60 distribution

Sharding Patterns

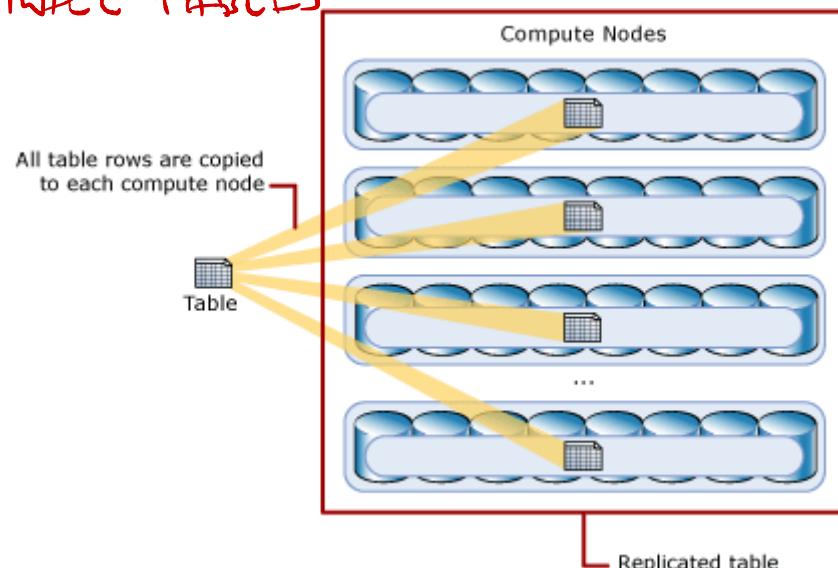


Replicated Tables

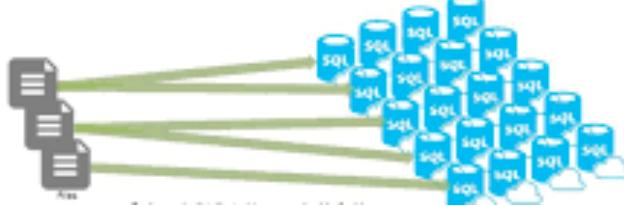
FAST QUERY PERFORMANCE FOR SMALL TABLES

- Caches a full copy on each compute node.
- Used for small tables

```
CREATE TABLE [dbo].[BusinessHierarchies](
    [BookId] [nvarchar](250) ,
    [Division] [nvarchar](100) ,
    [Cluster] [nvarchar](100) ,
    [Desk] [nvarchar](100) ,
    [Book] [nvarchar](100) ,
    [Volcker] [nvarchar](100) ,
    [Region] [nvarchar](100)
)
WITH
(
    CLUSTERED COLUMNSTORE INDEX
    DISTRIBUTION = REPLICATE
)
```



Round Robin tables



- Generally use to **load staging tables**
- **Distribute data evenly** across the table without additional optimization
- **Joins are slow**, because it requires to reshuffle data
- **Default** distribution type

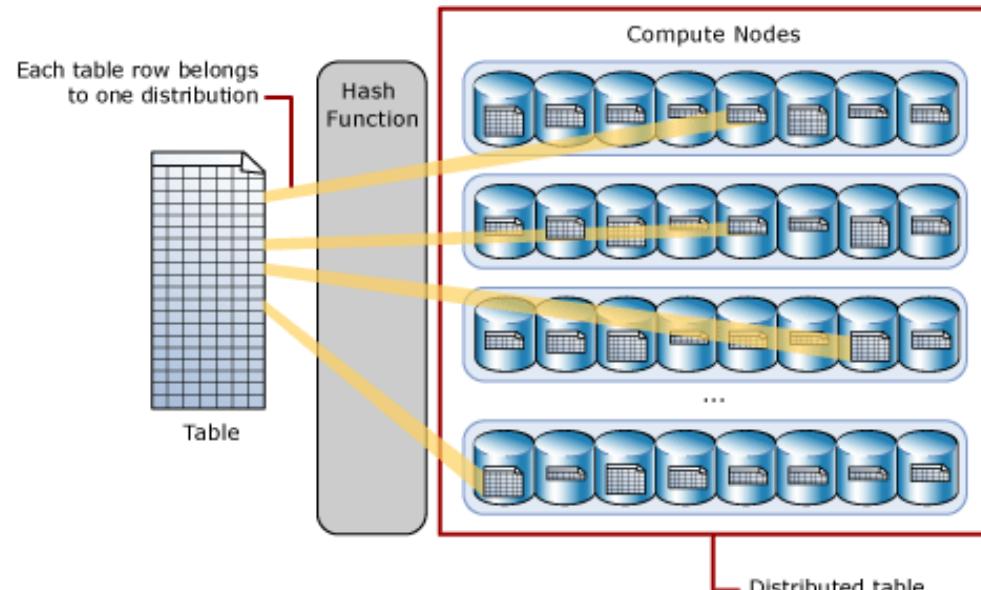
```
CREATE TABLE [dbo].[Dates](
    [Date] [datetime2](3) ,
    [DateKey] [decimal](38, 0) ,
    .
    .
    [WeekDay] [nvarchar](100) ,
    [Day Of Month] [decimal](38, 0)
)

WITH
(
    CLUSTERED COLUMNSTORE INDEX
    DISTRIBUTION = ROUND_ROBIN
)
;
```

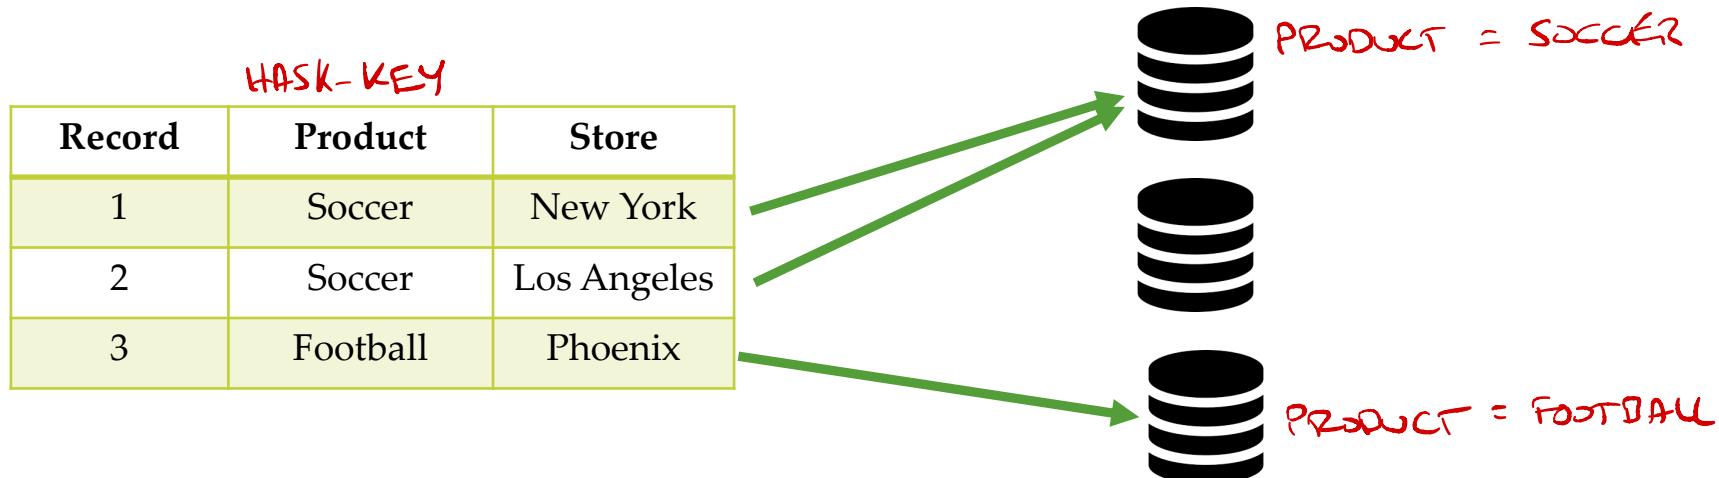
Hash Distribution Tables

- Highest performance for **large tables**
- Each row belong to one particular distribution **BASED ON HASH - (EM)**
- It is used mostly for larger tables

```
CREATE TABLE [dbo].[EquityTimeSeriesData](
    [Date] [varchar](30),
    [BookId] [decimal](38, 0),
    [P&L] [decimal](31, 7),
    [VaRLower] [decimal](31, 7)
)
WITH
(
    CLUSTERED COLUMNSTORE INDEX
    , DISTRIBUTION = HASH([P&L])
)
;
```



Hash Distribution Tables



Avoid Data Skew

PICK EVENLY DISTRIBUTED HASH KEYS TO AVOID SKEW



Even Distribution

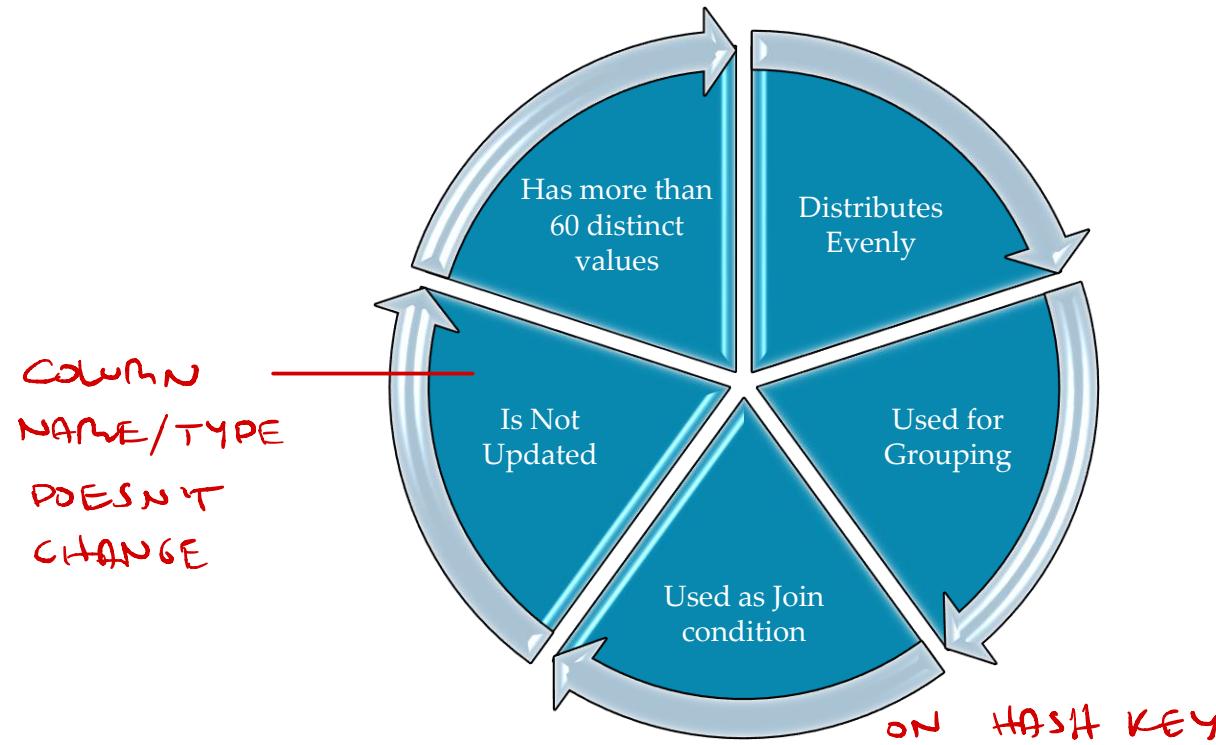


Distribution key

Determines the method in which Azure SQL Data Warehouse spreads the data across multiple nodes.

Azure SQL Data Warehouse uses up to 60 distributions when loading data into the system. \ DISTRIBUTIONS ARE ASSIGNED TO DIFFERENT COMPUTE NODES

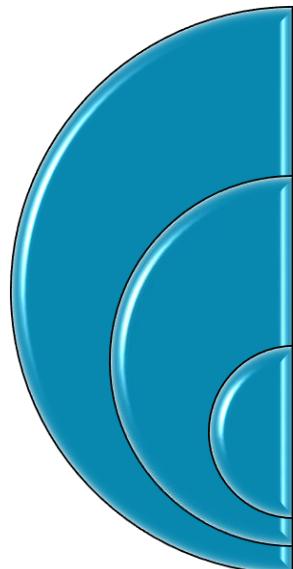
Good Hash Key



What Data Distribution to Use?

Type	Great fit for	Watch out if...
Replicated	Small-dimension tables in a star schema with less than 2GB of storage after compression	<ul style="list-style-type: none">Many write transaction are on the table (insert/update/delete)You change DWU provisioning frequentlyYou use only 2-3 columns, but your table has many columnsYou index a replicated table
Round-robin (default)	<ul style="list-style-type: none">Temporary/Staging tableNo obvious joining key or good candidate column.	Performance is slow due to data movement
hash	<ul style="list-style-type: none">Fact tablesLarge dimension tables	The distribution key can't be updated

Data types



Use the smallest data type which will support your data

Avoid defining all character columns to a large default length

Define columns as VARCHAR rather than NVARCHAR if you don't need Unicode

NVARCHAR USES 2X SPACE (BYTES) AS VARCHAR

Data types



The goal is to not only save space but also move data as efficiently as possible.

Data types



Some complex data types (XML, geography, etc)
are not supported on Azure SQL Data
Warehouse yet.

Table types

Clustered
columnstore

- Updateable primary storage method
- Great for read-only

Heap

- Data is not in any particular order.
- Use when data has no natural order.

Clustered Index

- An index that is physically stored in the same order as the data being indexed

Default table type

High compression
ratio

Clustered
columnstore

Ideally segments of
1M rows

No Secondary
Indexes

NO
ORDERING

No index on the data

Fast Load

Heap

No compression

Allows secondary indexes

CLUSTERED - INDEX

Sorted index on
the data

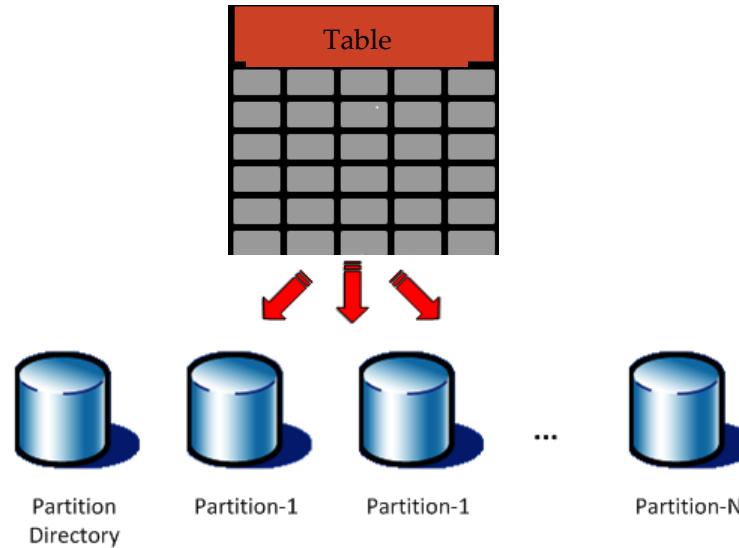
Fast singleton
lookup

Clustered B-Tree

No compression

Allows secondary
indexes

Table Partitioning





Partitioning

Table partitions enable you to divide your data into smaller groups of data

Improve the efficiency and performance of loading data by use of partition deletion, switching and merging

Usually data is partitioned on a date column tied to when the data is loaded into the database

Can also be used to improve query performance

Why Partitioning?

Easy load or
unload data

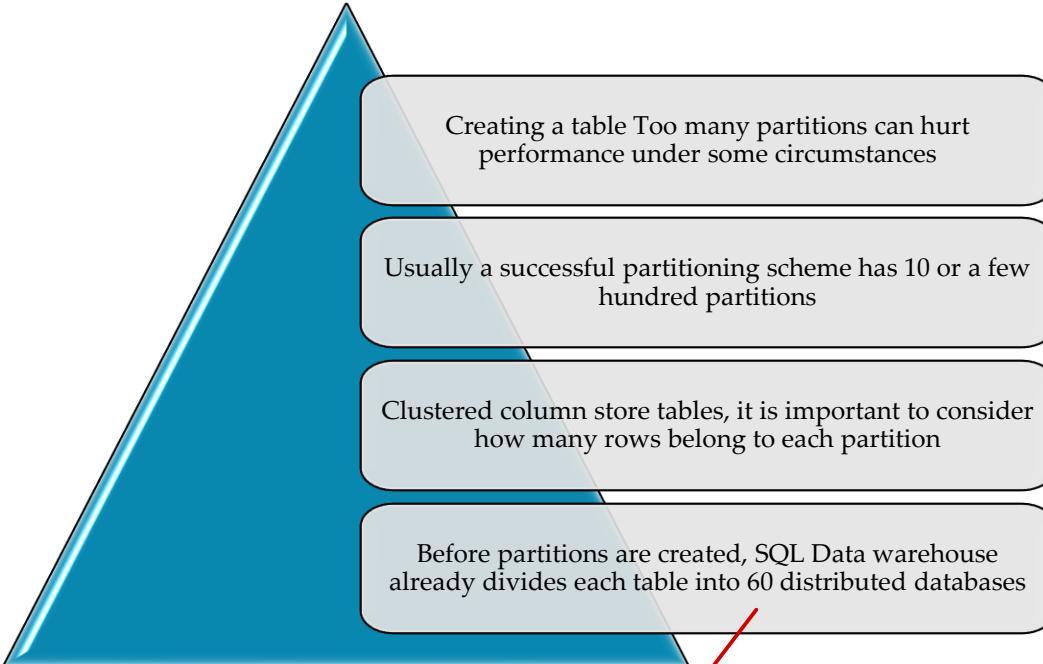
Easy
maintenance,
rebuilds or
reorganizes

Improve Query
Performance

BY TARGETING
SPECIFIC PARTITIONS

AVOIDS
FULL TABLE
SCANS

Partitions best practices



Creating a table Too many partitions can hurt performance under some circumstances

Usually a successful partitioning scheme has 10 or a few hundred partitions

Clustered column store tables, it is important to consider how many rows belong to each partition

Before partitions are created, SQL Data warehouse already divides each table into 60 distributed databases

BE MINDFUL NOT TO
OVER PARTITION



A highly granular partitioning scheme can work in SQL Server but hurt performance in Azure SQL Data Warehouse.

Example

OVER-PARTITIONING

60 Distributions



DAILY

365 Partitions



21900 Data Buckets

21900 Data
Buckets



Ideal Segment
Size (1M Rows)



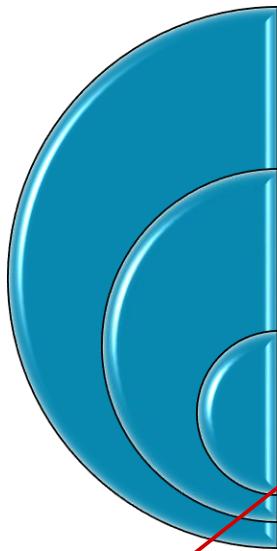
21 900 000 000 Rows



Lower Granularity (~~week, month~~)
can perform better depending on
how much data you have.

How do we apply these
principles to a Dimensional
model?

Fact Tables



Large ones are better as Columnstores

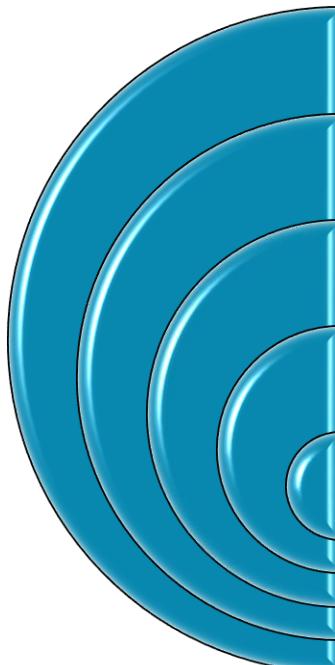
Distributed through Hash key as much as possible as long as it is even

Partitioned only if the table is large enough to fill up each segment

IDEAL SEGMENT (PARTITION) = 1m Rows

IDEAL TABLE SIZE = 60 PARTITION X 1m Rows = 60m Rows

Dimension Tables



Can be Hash distributed or Round-Robin if there is no clear candidate join key

Columnstore for large dimensions

Heap or Clustered Index for small dimensions

Add secondary indexes for alternate join columns

Partitioning not recommended

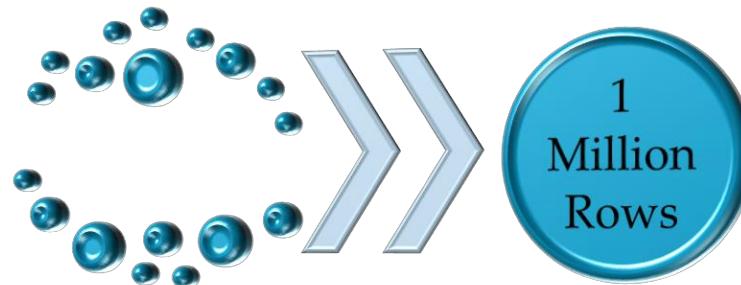
Best Practices for Data Laod

Data Warehouse Readers

Your DWUs have a direct impact on how fast you can load data in parallel

Optimize Insert Batch Size

Avoid trickle insert pattern. Ideal batch size is 1 million or more direct or in a file.

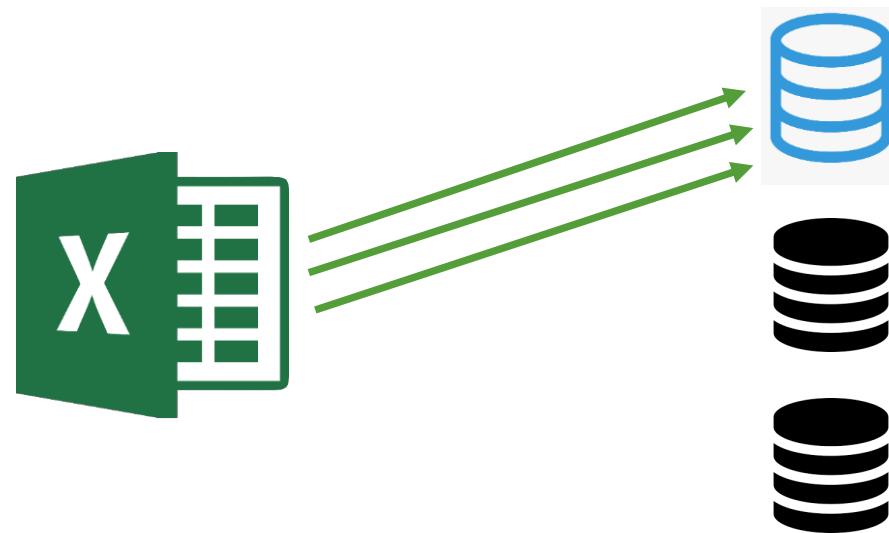


Combine rows to
make it a batch of
1 million

Ideal batch size

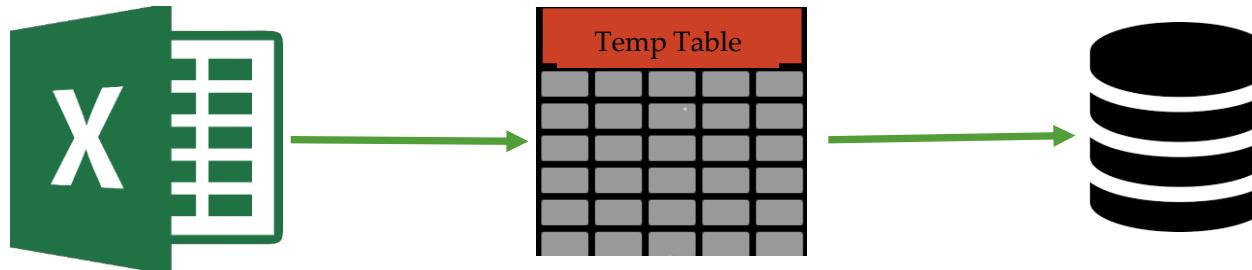
Avoid ordered data

Data ordered by distribution key can introduce hot spots that slow down the load operation



Using temporary tables

Stage and transform on a Temp Heap table before moving to permanent storage.



CREATE TABLE AS

```
CREATE TABLE #tmp_fct
WITH
(
DISTRIBUTION = ROUND_ROBIN
)
AS
SELECT *
FROM
[dbo].[FactInternetSales];
```

- Fully Parallel operation
- It is minimally logged
- It can change: distribution, table type, partitioning

Loading Methods

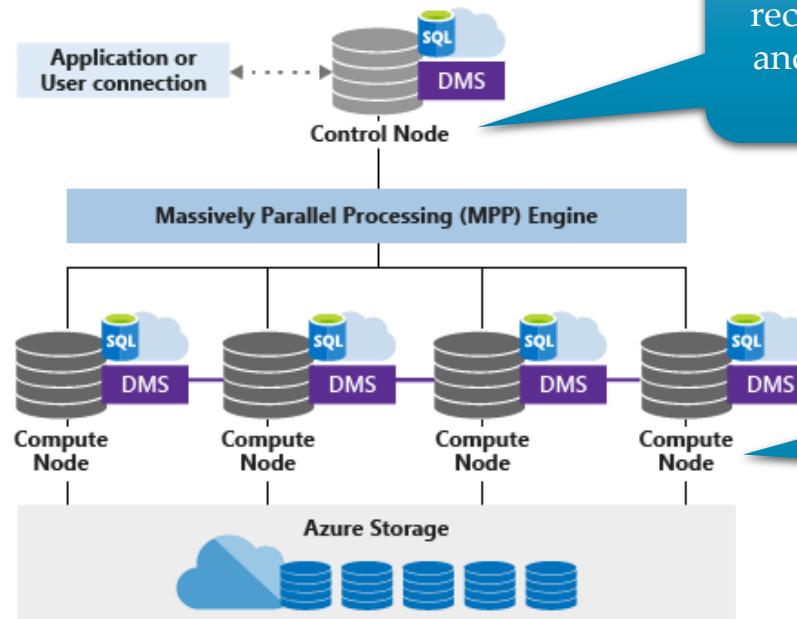
Single client loading methods

- SSIS
- Azure Data Factory
- BCP
- Can add some parallel capabilities but are bottlenecked at the control node

Parallel readers loading methods

- PolyBase
- Reads from Azure blob Storage and loads the contents into Azure SQL DW
- Bypasses the Control Node and loads directly into the Compute Nodes

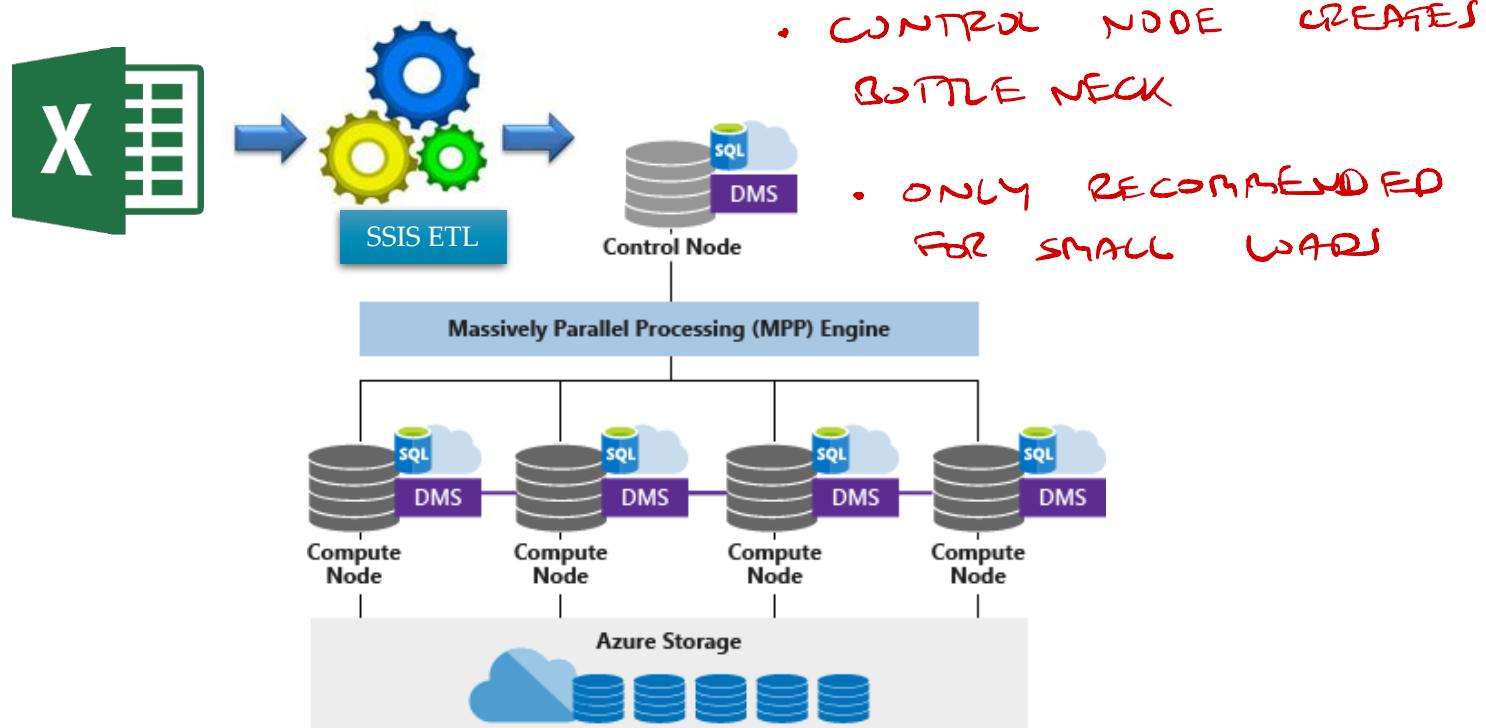
Control Node



The **Control** node receives connections and orchestrates the queries

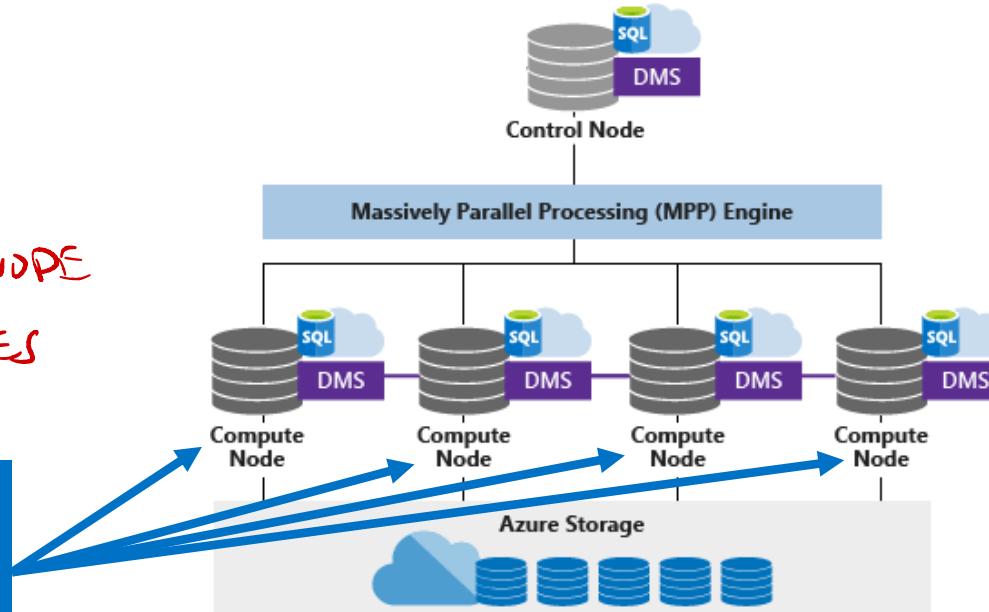
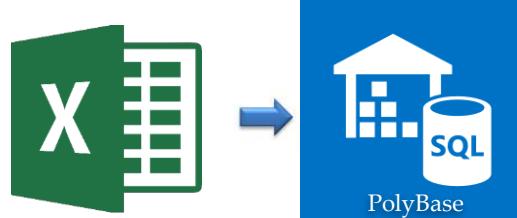
The **Compute** nodes do processing on the data and scale with the DWUs.

Loading with SSIS



Loading with PolyBase

- RECOMMENDED FOR
LARGE LOADS
- BYPASSED CONTROL NODE
FOR MPP CAPABILITIES





PIPE, COMMA, TAB
DELIMITED

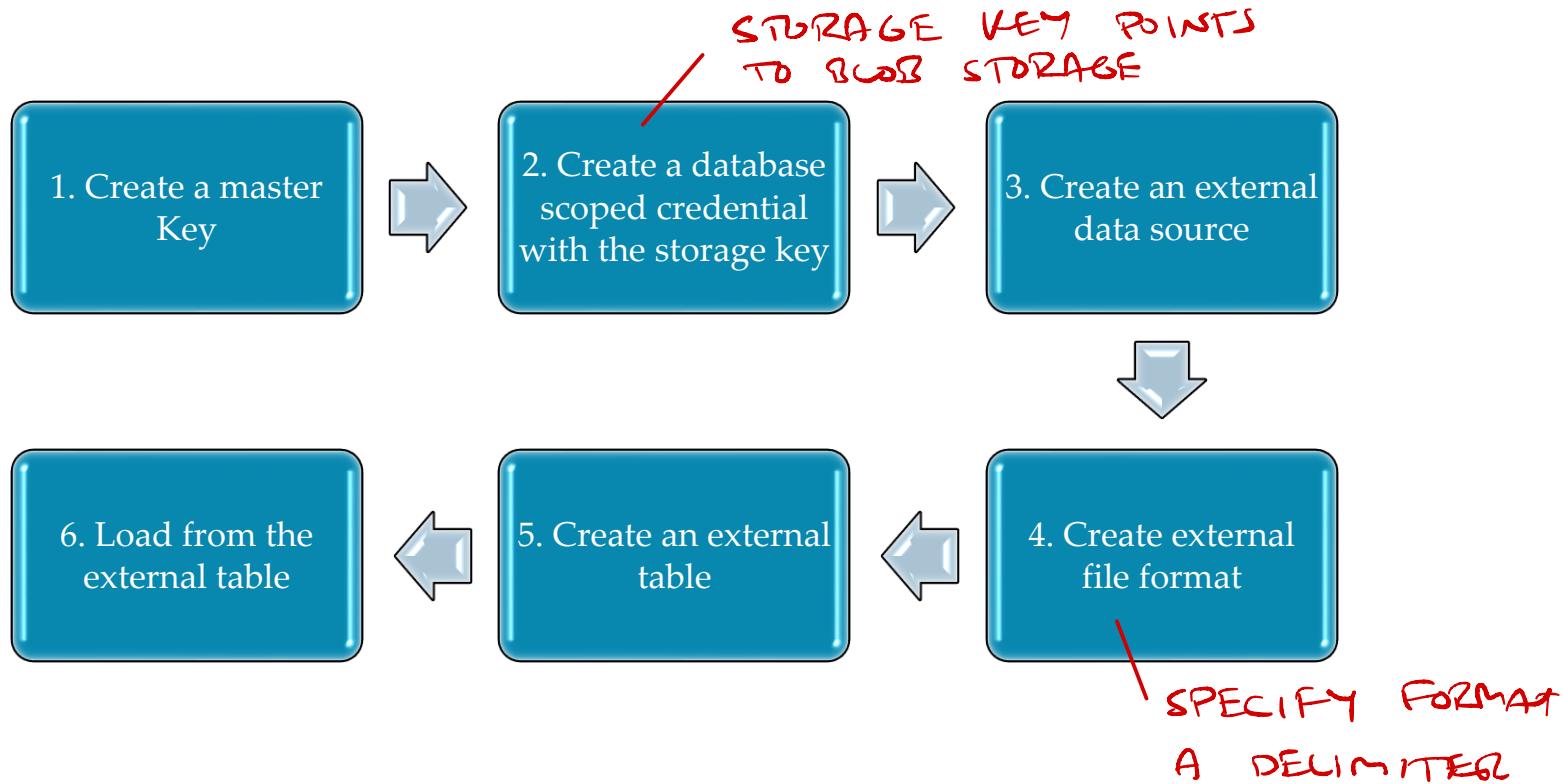
PolyBase can load data from UTF-8
delimited text files and popular
Hadoop file formats like RC File,
ORC, and Parquet

CAN COMPRESS FILES BEFORE LOADING



Multiple readers will not work against a
compressed file

PolyBase Setup



Demo – PolyBase

1. Export table to flat file
2. Create blob storage account
3. Upload flat file to blob storage
4. Run PolyBase 6 steps process
5. Monitor and confirm successful migration
6. Confirm 60 distributions in destination table



Azure
Synapse
Analytics

Azure SQL Warehouse Scaling

Scaling in SQL Data Warehouse

- SQL DW allows us to **scale or pause at will**
- Scale up during heavy demand
- Pause to cut cost
- **DWUs are CPU, memory, and I/O bundled into units of compute scale**
- Increasing DWU's
 - increase query performance
 - Also increase maximum number of concurrent queries and concurrent slots
- **Modify with GUI, PowerShell, or TSQL**



Azure SQL DW Backup and Restore

DIFFERENT FROM SQL DATABASE

- Snapshots of your data warehouse are taken throughout the day creating restore points
- These restore points are available for 7 days
 - Retention period of 7 days cannot be changed
- SQL pool supports an 8 hour recovery point objective (RPO)
- Replicated to paired region once a day
- You can also take user-defined snapshots
 - Retention period 7 days cannot be changed
 - 42 max restore point possible
 - Can be created using PowerShell or portal
- When you drop a SQL pool, a final snapshot is created and saved for seven days
 - SQL Pool should not be in paused state *



CAN RESTORE FROM AUTOMATIC OR USER-DEFINED RESTORE POINTS

Azure SQL DB vs. Azure SQL DW

CRUD = CREATE, READ, UPDATE, DELETE

OLTP/CRUD

SYMMETRIC
MULTI-PROCESSING

SMP

Vertical scale

No Polybase



Azure
SQL Database



Polybase

Azure
SQL Data Warehouse

OLAP/querying and reporting

GO COMPUTE NODES

MPP

Horizontal Scale

Can pause the
virtual server
to save cost

Dynamic Data Masking

CAN IMPLEMENT IN AZURE PORTAL : SQL DB > DYNAMIC DATA MASKING

ID	PersonName	EmailAddress	CreditCardNumber	SocialSecurityNumber
1	Anoop Kumar	abcdefg@hotmail.com	1234-5678-4321-8765	123-45-6789
1	Rahul Gupta	amitguptaabcfgh@hotmail.com	8765-1234-5678-4321	231-45-6787
1	Amit Goel	amitgoelabcdefg@hotmail.com	4321-1234-5678-4321	321-45-6700

OR WITH T-SQL
SCRIPTS

ID	PersonName	EmailAddress	CreditCardNumber	SocialSecurityNumber
9590	AXXXr	aXXX@XXXX.com	xxxx-xxxx-xxxx-8765	xxxx
7604	RXXXa	aXXX@XXXX.com	xxxx-xxxx-xxxx-4321	xxxx
8453	AXXXI	aXXX@XXXX.com	xxxx-xxxx-xxxx-4321	xxxx

Random Number function – generates random number based on selected boundaries and actual data type. Can be applied only numbers, not string

Email function – Exposes the first letter and replace the domain with xxx.com

Default function – Full masking of data. For numeric – 0 For String – XXXX characters

Custom String function – You can define the exposed prefix, the padding string and exposed suffix

Credit Card function - Only the last four digits of Credit card are shown



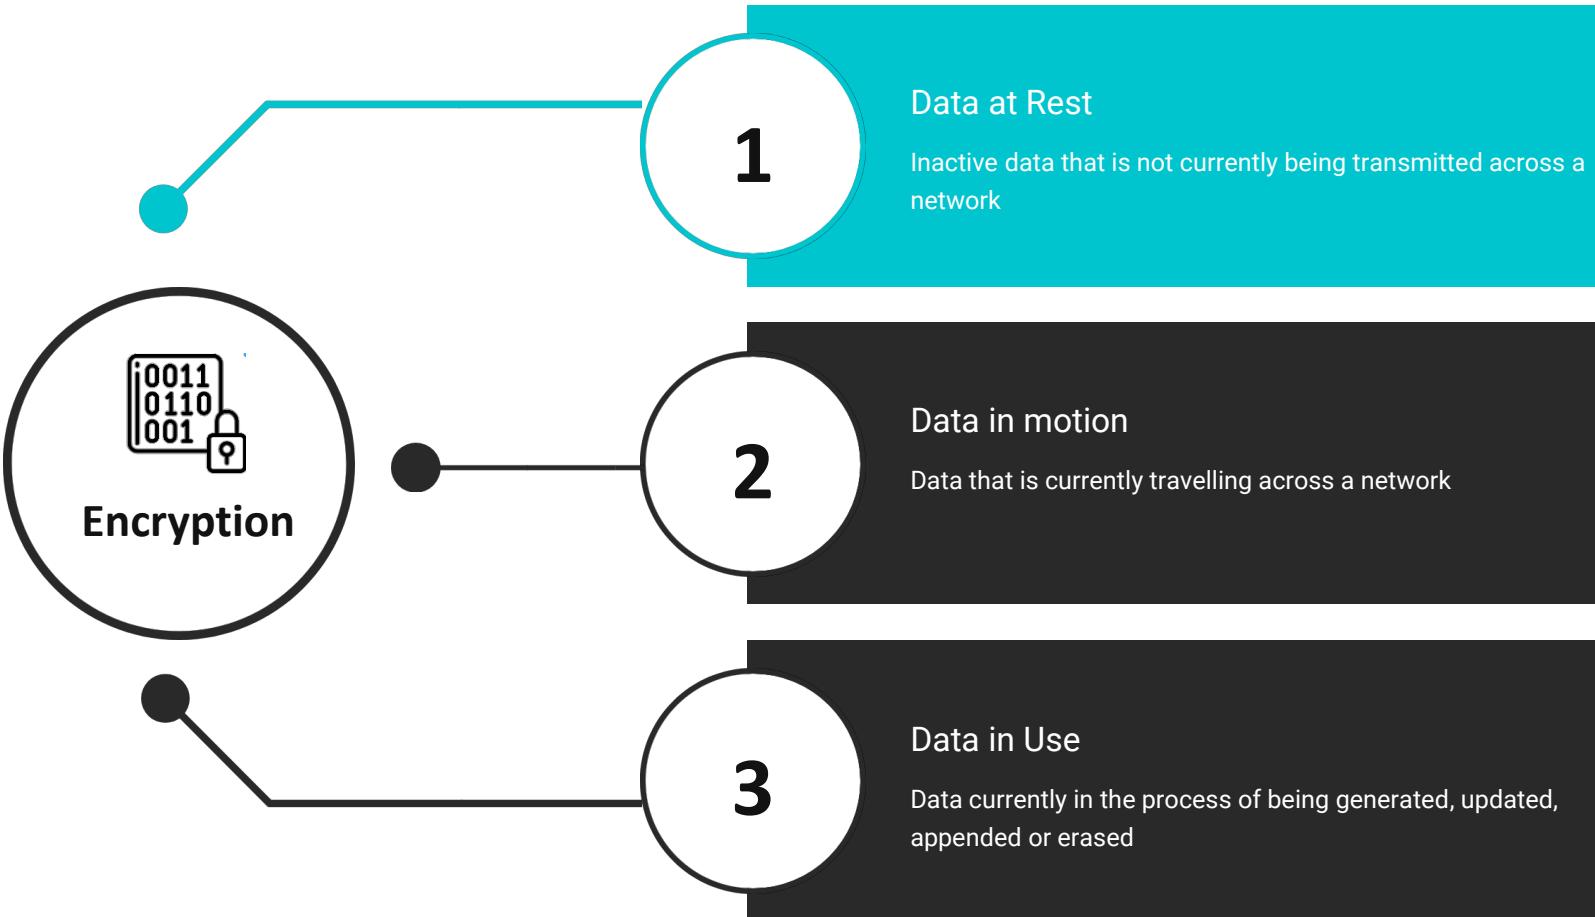
Dynamic Data Masking

- **Limit the exposure of sensitive data to non-privileged users**
 - You can decide the level of exposure of data
- **No change in physical layer**
 - Data in the database is not changed
 - Not the same as data encryption
- **No additional development effort needed at application level**
- **Security: Should not be used as a primary security layer**
 - Dynamic Data Masking should not be used as an isolated measure to fully secure sensitive data
 - ad-hoc query permissions can apply techniques to gain access to the actual data.
- **Other considerations**
 - Masked columns can be updated if user has permission
 - Export masked from source data results in masked data in target table

Dynamic Data Masking



- How to Enable DDM?
 - Portal **OVERVIEW > DYNAMIC DATA MASKING**
 - Powershell
 - Get-AzSqlDatabaseDataMaskingRule
 - New-AzSqlDatabaseDataMaskingRule
 - Remove-AzSqlDatabaseDataMaskingRule
 - Set-AzSqlDatabaseDataMaskingRule
 - Rest API



Types of Encryption



SAME KEY TO ENCRYPT
AND DECRYPT



Highly performant

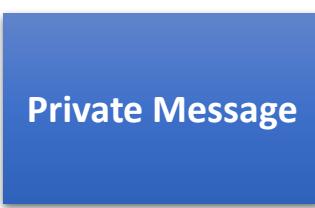


Secure key handling is difficult

Types of Encryption

Asymmetric

ENCRYPT



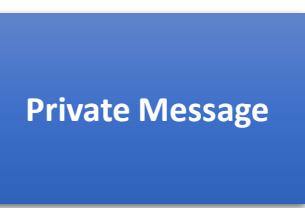
Public Key

(^R(^*()os*^s

DECRYPT



Private Key



DIFFERENT KEYS TO
ENCRYPT & DECRYPT



Set of keys (Public and Private) where one is used to encrypt and other to decrypt

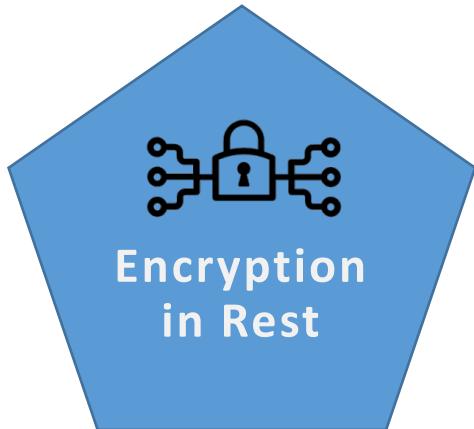


Computationally more expensive, but solves the key handling problem

PUBLIC KEYS CAN BE STORED & SHARED IN OWN KEY VAULT

Encryption in Rest

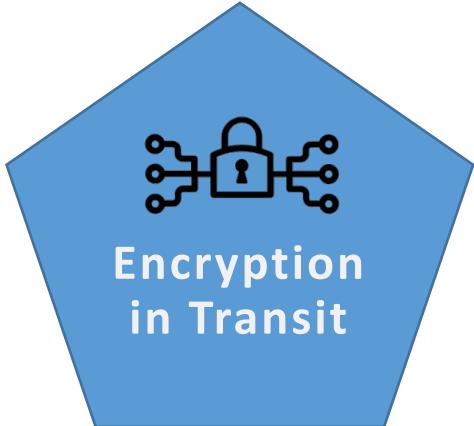
OPTION TO TURN OFF IN SQLDB & SYNAPSE
SQLDB SETTINGS CAN BE SET AT THE SERVER LEVEL



- By default all storage services encrypted data at rest
- 256-bit AES encryption
- Keys are stored in Azure Key Vault

OPTION TO USE MSFT MANAGED (DEFAULT)
KEY OR CUSTOMER-MANAGED KEY

Encryption in transit



- Any communication over the internet to Azure services is typically encrypted via SSL/TLS protocols
- Utilize site-to-site VPN or point-to-site VPN connections
WHEN USING ON-PREM TO VPN
- Or utilize ExpressRoute
- ExpressRoute communication is a private connection and not encrypted but workloads can be encrypted via SSL/TLS
- Storage services can be configured to require secure transfer
UNDER CONFIGURATION PAGE

Types of Encryption



Deterministic encryption:

- This will always generate the same encrypted value for any plain text value.
- You can perform point lookups, equality joins, grouping and indexing on encrypted columns.

Randomized encryption:

- This is more secure than deterministic encryption because the encrypted value is generated in a less predictable manner.
- But you can't perform searching, grouping, indexing or joins on encrypted columns.

Types of Encryption



Column Encryption Key (CEK)

- Used to encrypt values in specific columns
- Encrypted versions of each CEK is stored in the database.

Column Master Key (CMK)

- Used to encrypt all the CEKs
- Must be stored externally in a secure key store
- Key store providers: Azure key vault, certificate store, HSM